


A Little Bit More of Application Layer



Data distribution
methods


Content
Distribution
Networks

Peer-to-Peer

Internet Video

Data Distribution

- ◆ Client-server model
 - HTTP, NFS, AFP, SCP, RSYNC, ...

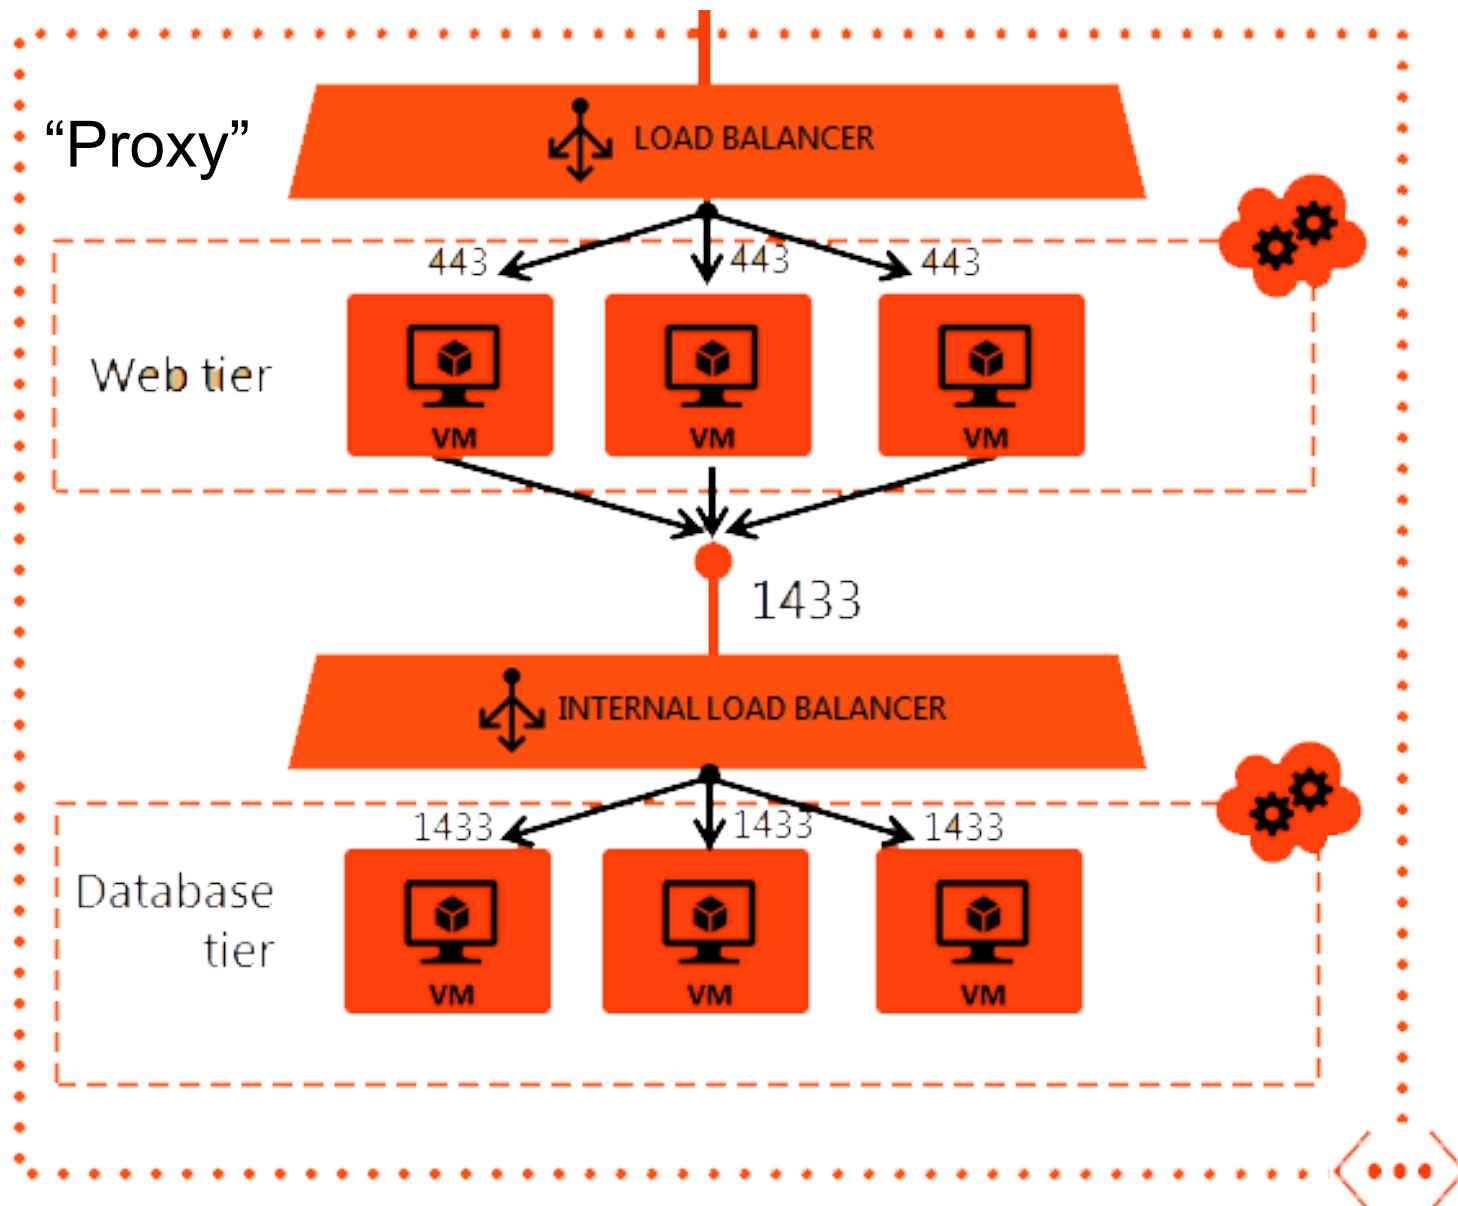


Main concern: how to scale?

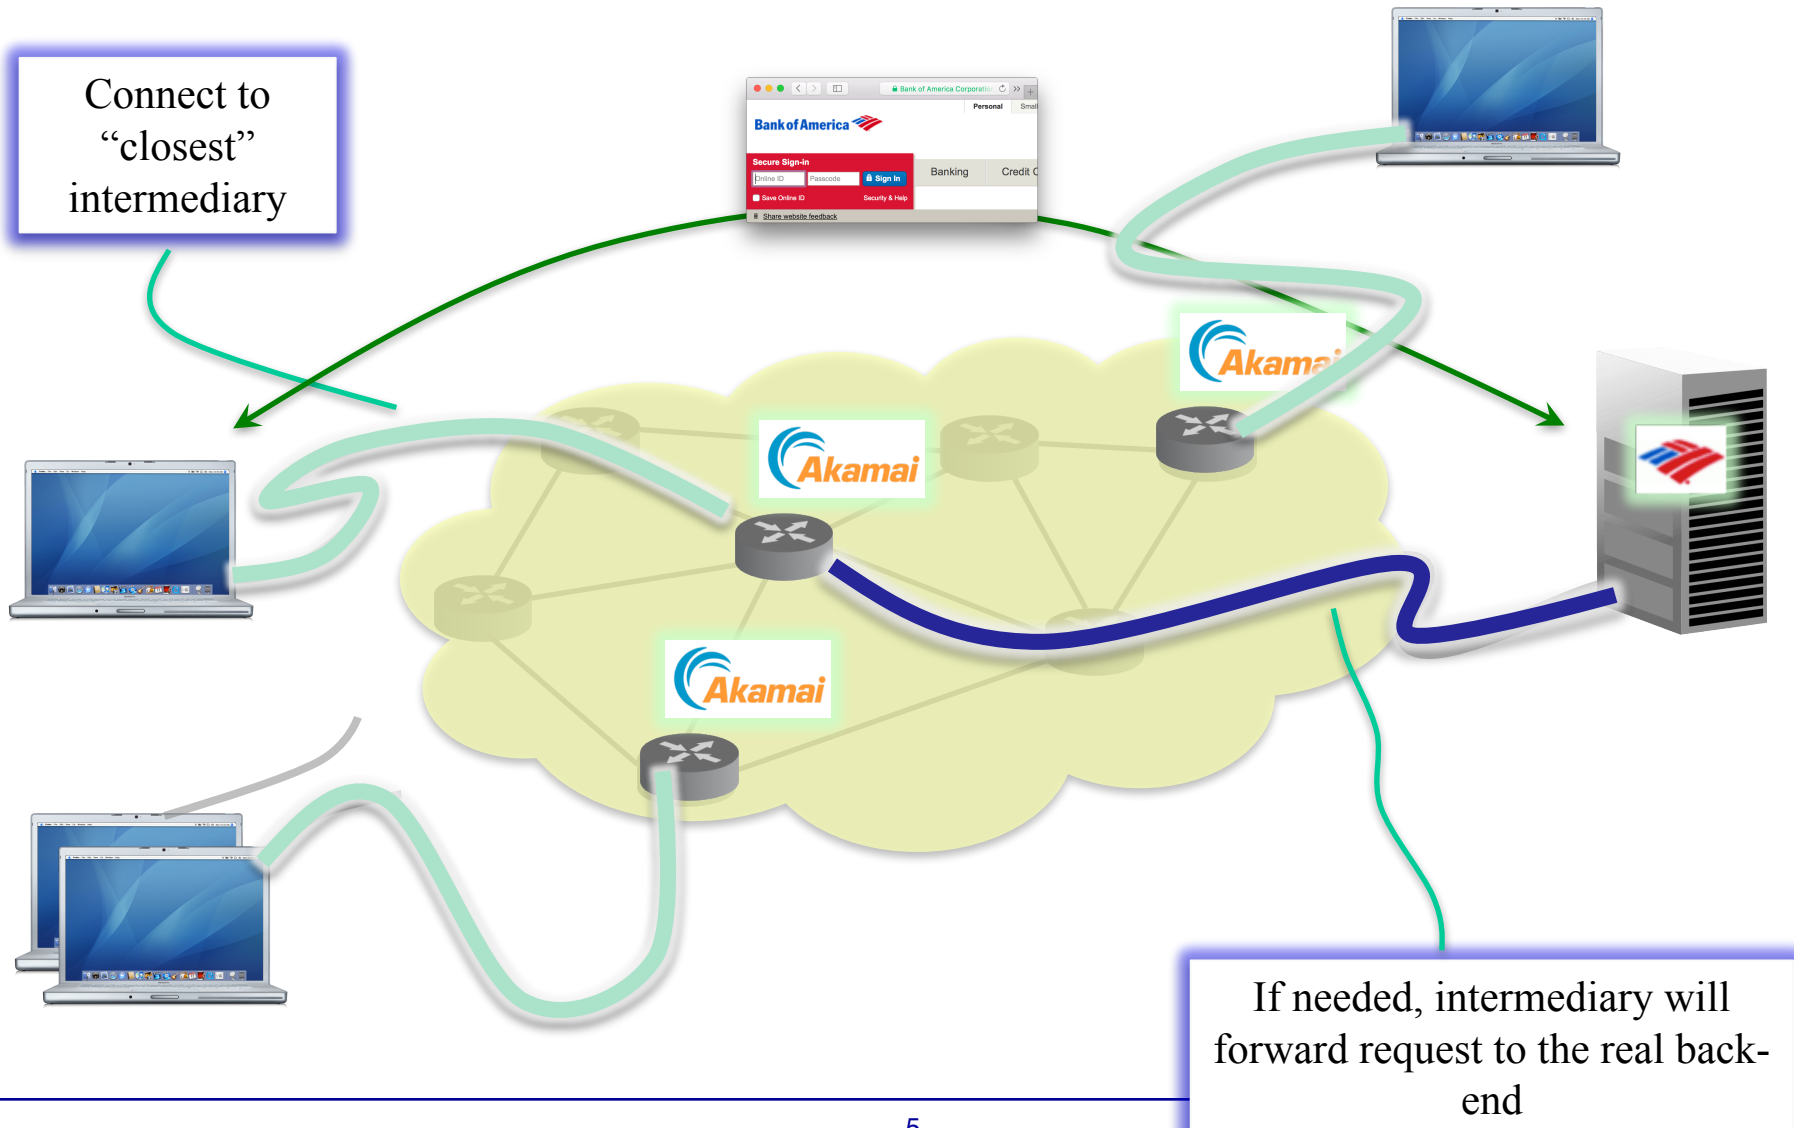
Scaling Data Distribution

- ◆ Remember DNS from previous lecture
- ◆ Does it scale?
- ◆ How?
 - Replication of authoritative servers
 - Aggregation and suppressing of similar requests by caching resolvers

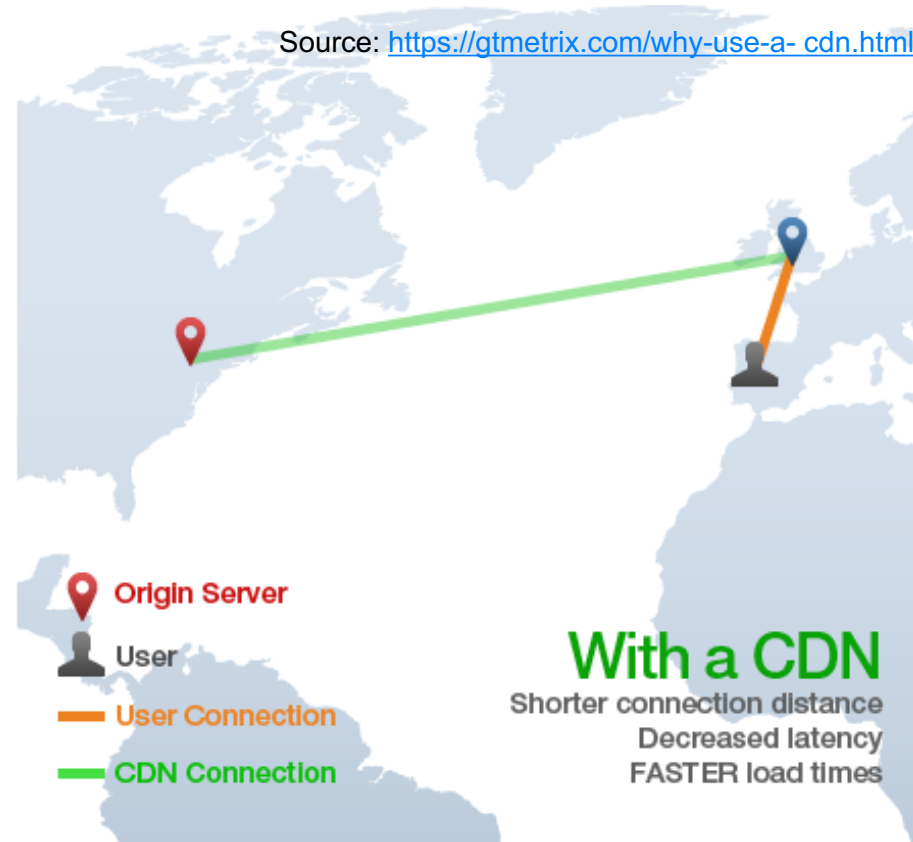
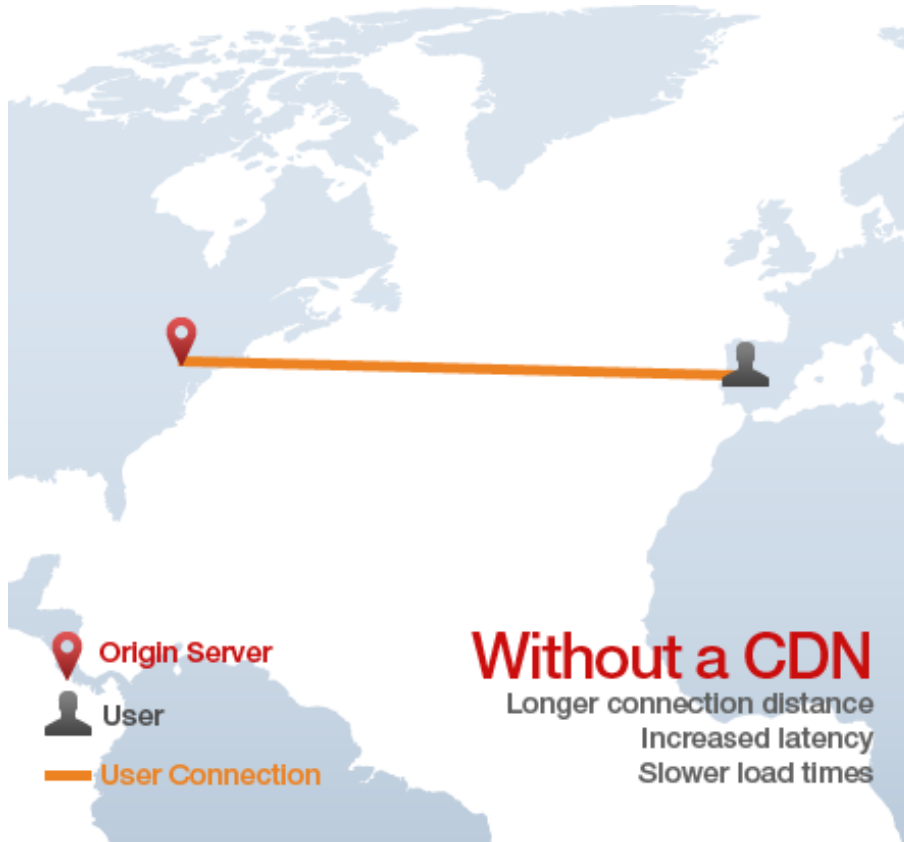
Scaling Data Distribution: Load Balancers



Scaling Data Distribution: Content Delivery Networks (CDNs)



With And Without CDNs



CDN Vendors

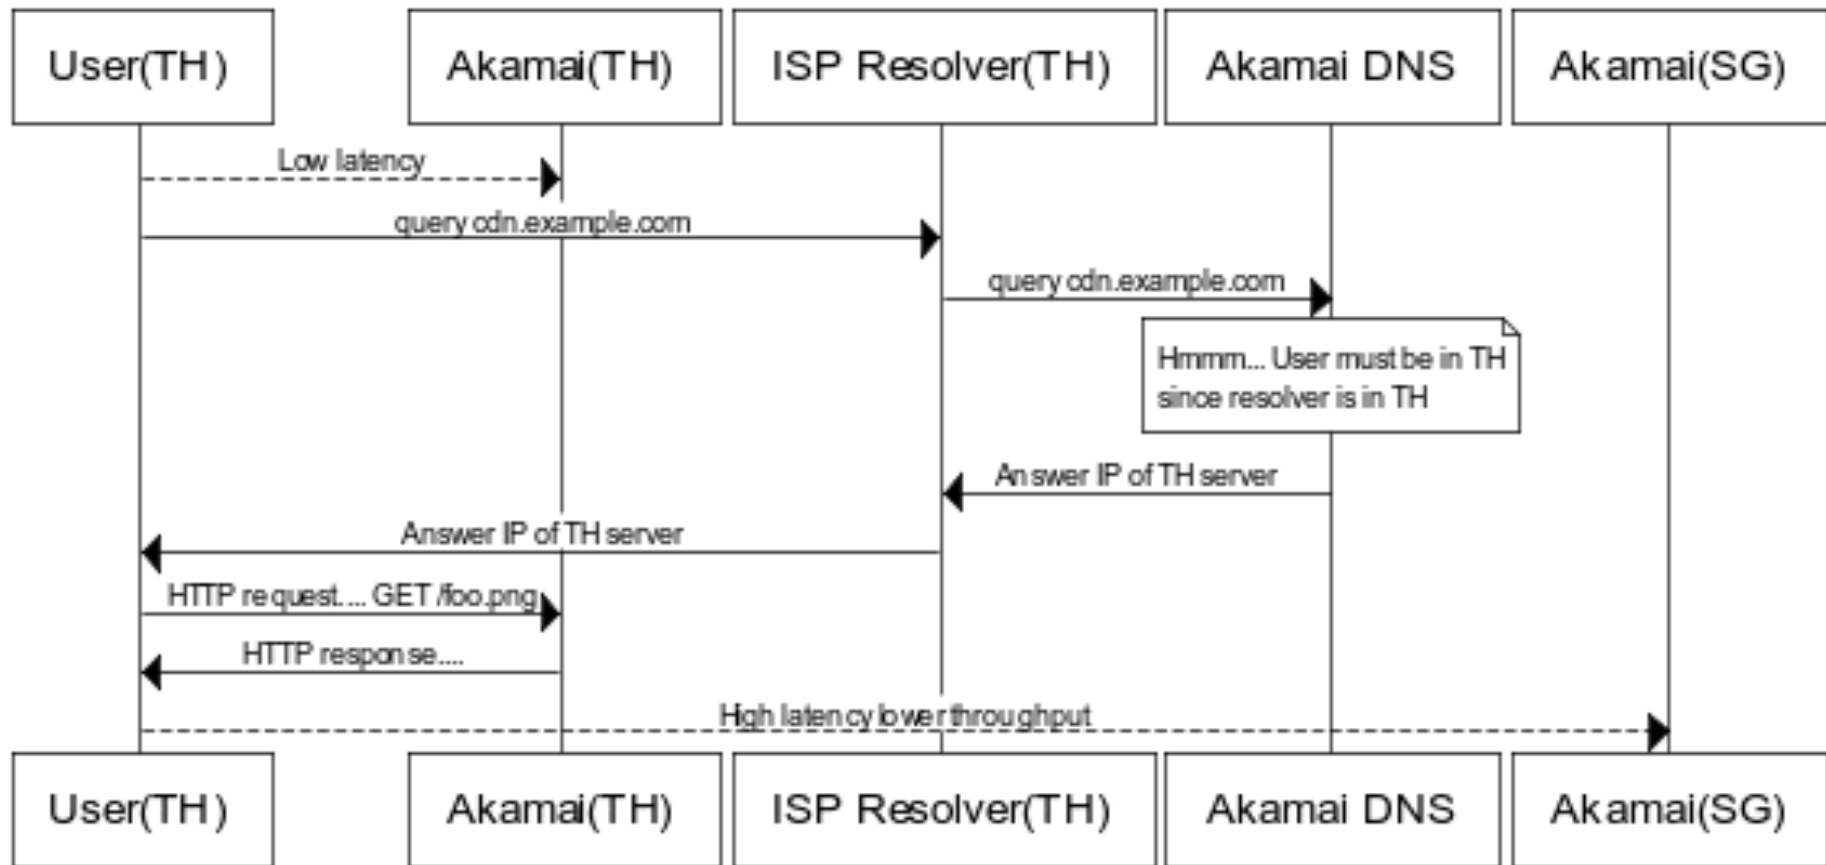
- ♦ [Akamai](#)
- ♦ [Alcatel Lucent](#) (carrier platform)
- ♦ [Allot Communications](#) (traffic management)
- ♦ [Amazon](#)
- ♦ [ARA Networks](#) (traffic management)
- ♦ [Aryaka](#)
- ♦ [Blue Coat](#) (transparent caching)
- ♦ [Broadpeak](#) (carrier platform)
- ♦ [BTI Systems](#) (traffic management)
- ♦ [CDNetworks](#)
- ♦ [Cedexis](#) (traffic management)
- ♦ [CDN77](#)
- ♦ [ChinaCache](#)
- ♦ [Cisco](#) (carrier platform)
- ♦ [Conversant](#) (carrier platform)
- ♦ [Comcast](#)
- ♦ [Conviva](#) (analytics)
- ♦ [DeepField](#) (analytics)
- ♦ [Edgeware](#) (carrier platform)
- ♦ [Ericsson](#) (carrier platform)
- ♦ [Fastly](#)
- ♦ [Fortinet](#) (traffic management)
- ♦ [Hibernia Networks](#)
- ♦ [Highwinds](#)
- ♦ [Huawei](#) (carrier platform/transparent caching)
- ♦ [Instart Logic](#)
- ♦ [Internap](#)
- ♦ [Jetstream](#) (licensed CDN)
- ♦ [Juniper](#) (transparent caching)
- ♦ [LeaseWeb](#)
- ♦ [Level 3](#)
- ♦ [Limelight Networks](#)
- ♦ [MaxCDN](#)
- ♦ [Microsoft](#) (Windows Azure)
- ♦ [MileWEB](#)
- ♦ [Mirror Image](#)
- ♦ [OnApp](#) (traffic management)
- ♦ [PeerApp](#) (transparent caching)
- ♦ [Qwilt](#) (transparent caching)
- ♦ [Radware](#)
- ♦ [Revsu](#) (mobile CDN)
- ♦ [Solbox](#) (licensed CDN)
- ♦ [Swiftserve](#) (licensed CDN)
- ♦ [Tata Communications](#)
- ♦ [Verizon EdgeCast](#)
- ♦ [Vidscale](#) (carrier platform)
- ♦ [Yottaa](#)

<http://blog.streamingmedia.com/2014/07/cdnvendors.html>

How to Connect to “Closest” CDN Node

- ◆ Abuse DNS
 - send DNS query for the name
 - DNS server
 - extracts IP address from the query
 - maps IP to “closest” network or geo coordinate
 - returns set of IP addresses that seem to be closer to you
- ◆ This is what DNS Scavenger hunt challenge is about

How CDNs work



www.websequencediagrams.com

Source: <http://www.cdnplanet.com/blog/which-cdns-support-edns-client-subnet/>

A Caveat


- ◆ DNS "redirection" works great if users use ISP-provided caching resolvers
 - Source IP of the caching resolver determines which answer authoritative DNS resolver returns
- ◆ But there is a problem with "public" DNS servers
 - Google's 8.8.8.8, 8.8.4.4
 - What answer to return to caching resolver 8.8.8.8?
 - CDN's DNS server will see request from google, not from client
- ◆ Solution
 - DNS protocol was extended to include "edns-client-subnet" option, which is set by caching resolver to tell authoritative name server IP of a client
 - How does this work with caching?

Other Ways to Scale Data Distribution

- ◆ Peer-to-peer networks
 - Unstructured p2p networks
 - BitTorrent
 - Gnutella, Gossip, and Kazaa
 - Structured p2p networks (will not cover, but recommend reading in the textbook)
 - Chord, CAN, Tapestry, Pastry, Kademlia

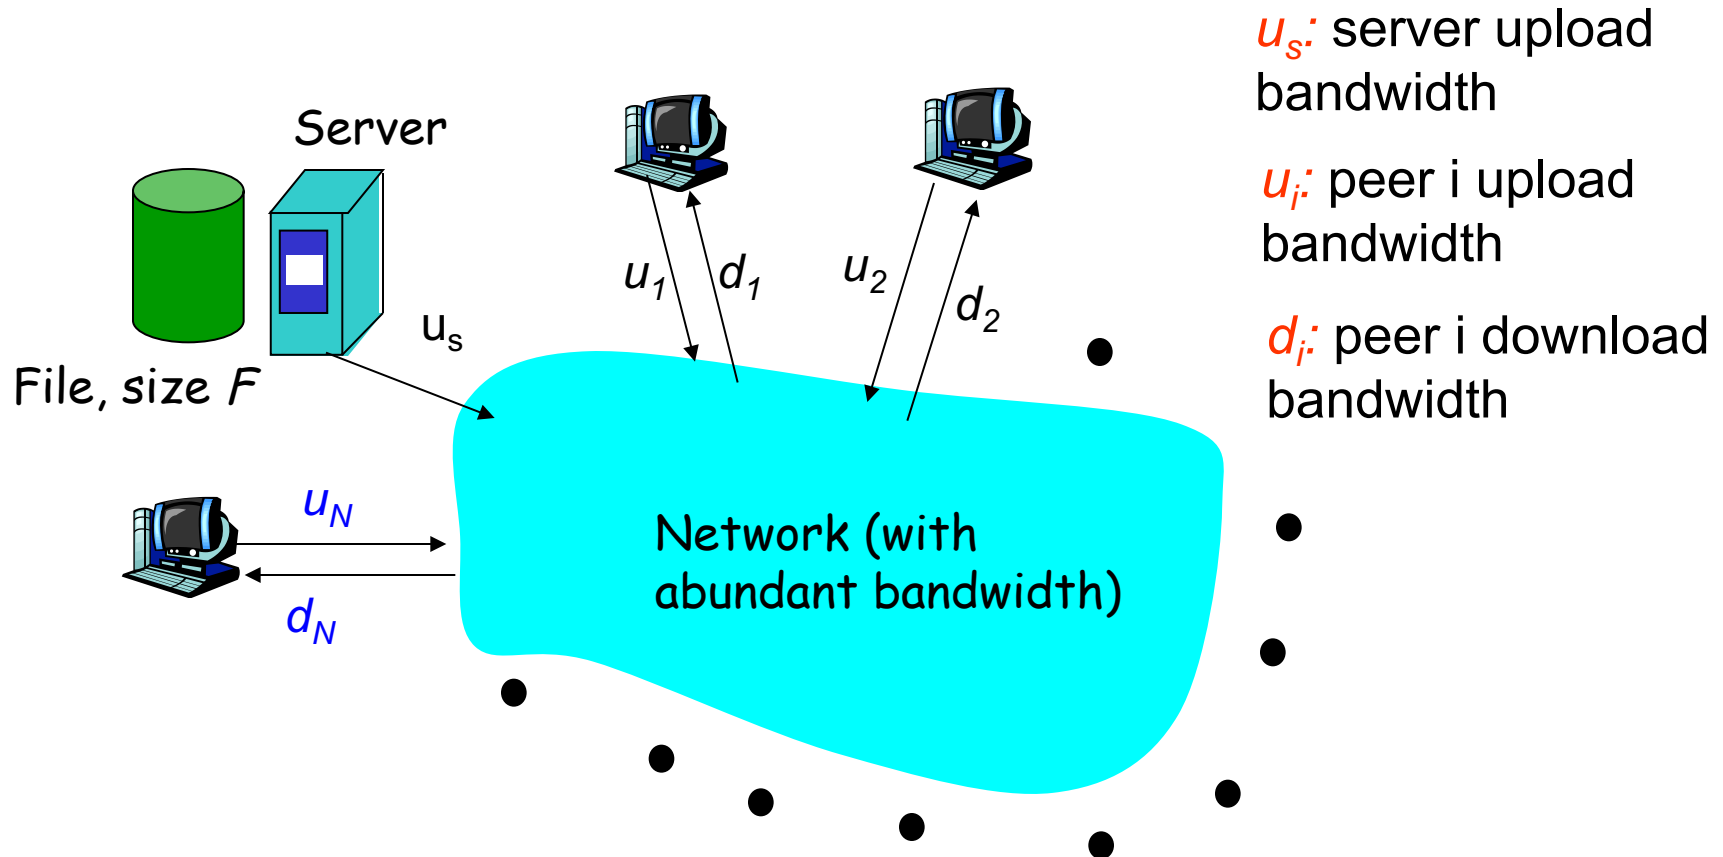
Properties of Peer-to-Peer Systems

- ◆ No “always-on” server
- ◆ Arbitrary end systems directly communicate
- ◆ Peers are intermittently connected and change IP addresses




File Distribution: Client-Server vs P2P

Question : How much time needed to distribute a file from one server to N clients?



File distribution time: client-server

- ♦ server sequentially sends (upload) N copies:
 - time to send one copy:
 - F/u_s
 - time to send N copies:
 - $N * F/u_s$ (lower bound)
- ♦ client i takes F/d_i time to download




Time to distribute F to N clients using Client-server approach $= d_{cs} = \max \{ N * F/u_s, F/\min_i(d_i) \}$

increases linearly with N

File distribution time: P2P

- ♦ server must send one copy:
 - Time to send one copy: F/u_s
- ♦ client i takes F/d_i time to download
- ♦ Total $N \cdot F$ bits must be downloaded (aggregate)
 - max upload rate is $u_s + \sum u_i$



*time to distribute F
to N clients using
P2P approach*


$$d_{p2p} \geq \max \{ F/u_s, F/\min(d_i), N \cdot F / (u_s + \sum u_i) \}$$

increases linearly in N ...

... but so does this item as each peer brings service capacity

Client-server vs. P2P: example

client upload rate = u , $F/u = 1$ hour, $u_s = 10u$, $d_{min} \geq u_s$



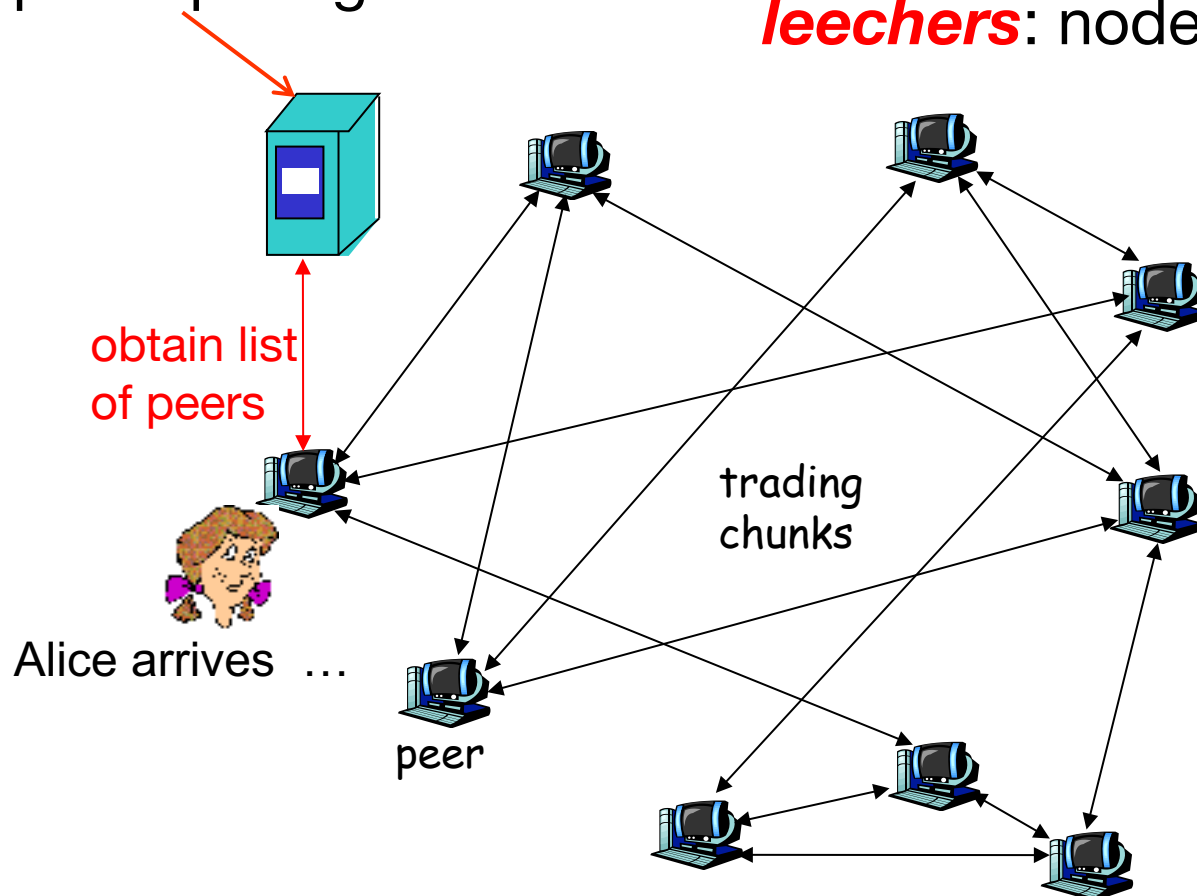
P2P File distribution: BitTorrent

tracker: centralized server that keeps track of participating nodes

torrent: group of peers exchanging chunks of a file

seeds: node with entire content

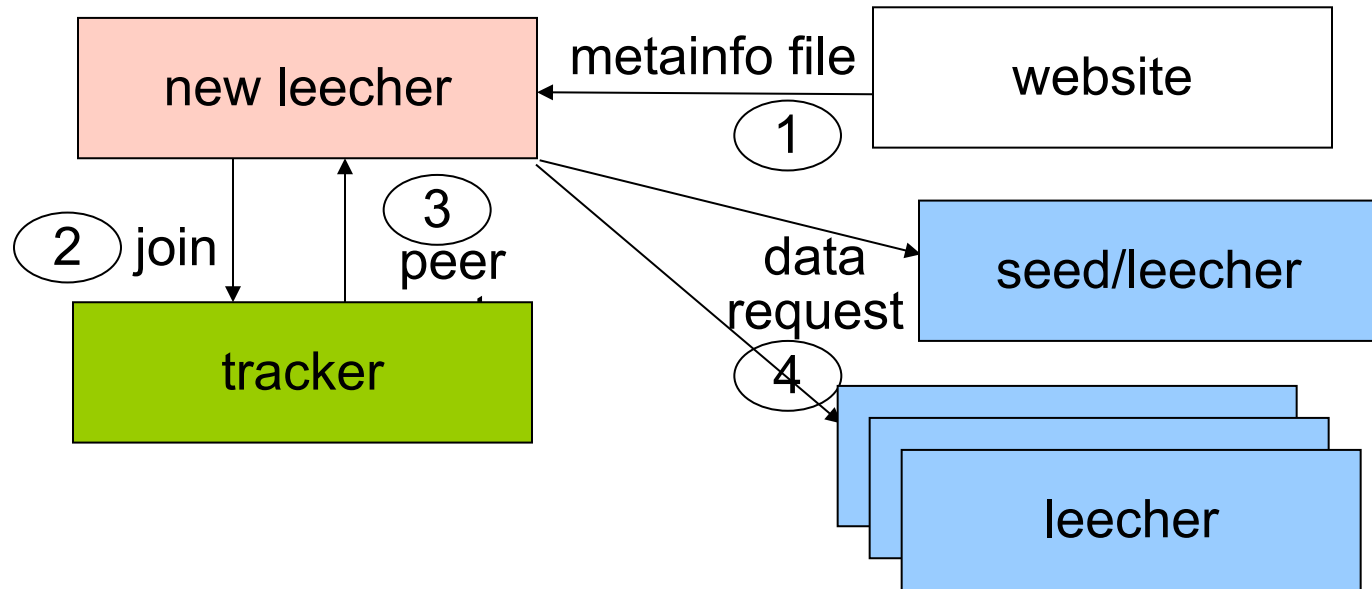
leechers: node still downloading



BitTorrent Philosophy

- ◆ Author: Bram Cohen
- ◆ Based on Tit-for-tat
- ◆ Incentive - Uploading while downloading
- ◆ Pieces of files

BitTorrent – joining a *torrent*



1. obtain the *metainfo file* from a website
2. contact the control server (tracker)
3. obtain a *peer set* (contains seeds & leechers)
4. contact peers in the peer set to request data

BitTorrent: moving data

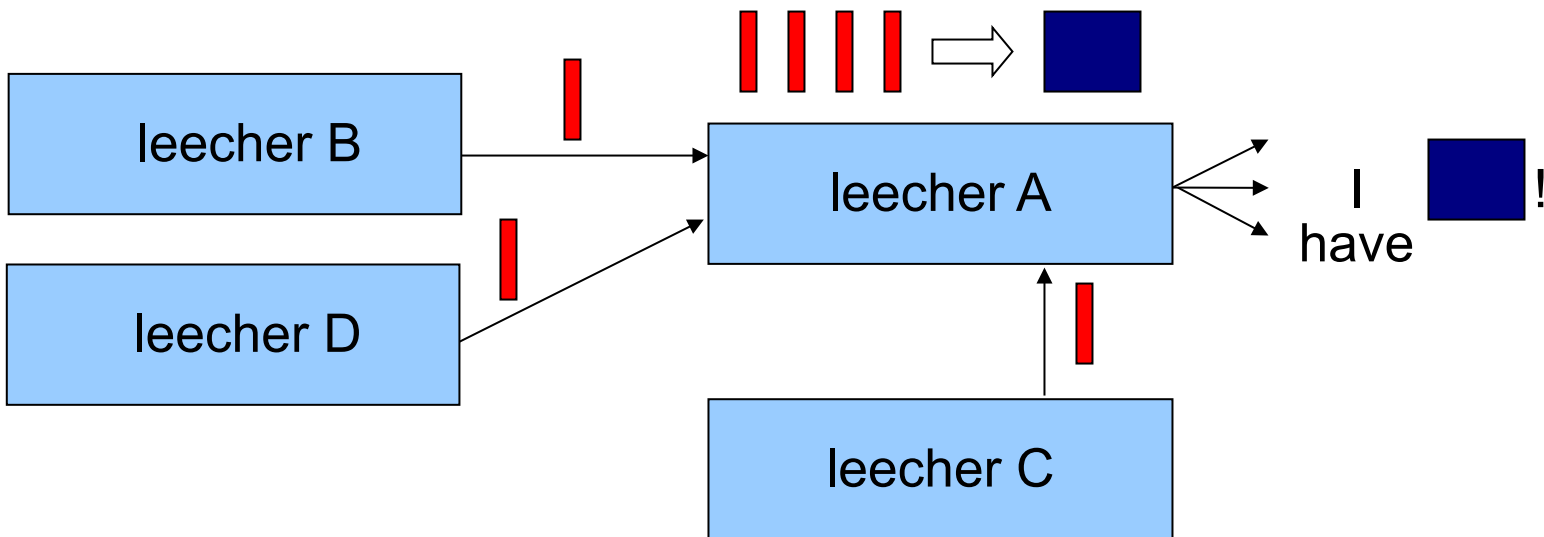
Pulling Chunks

- ◆ file divided into 256KB *chunks*
- ◆ at any given time, different peers may have different subsets of file chunks
- ◆ periodically, a peer (Alice) asks each neighbor for the list of chunks they have
- ◆ Alice sends requests for her missing chunks
 - rarest first


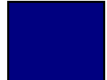
Sending Chunks: tit-for-tat

- ◆ Alice sends chunks to four neighbors currently sending her chunks *at the highest rate*
 - re-evaluate top 4 every 10 sec
- ◆ every 30 secs: randomly select another peer, starts sending chunks
 - “optimistically unchoke”
 - The newly chosen peer may become one of Alice’s top 4


BitTorrent – downloading data



After A learned peers B & C's chunk list:

- download the *rarest* missing piece first
- download data *blocks*  in parallel
- verify the *piece*  using hashes in the metainfo file
- advertise received pieces to the entire peer set

BitTorrent – incentives mechanism



- ◆ all leechers periodically calculate download rates from others, and only upload to the fastest (*regular unchoke*)
 - allows new leechers to get their first data pieces
- ◆ *optimistic unchoke*: periodically select a peer at random and send it data

BitTorrent Security

- ◆ Let's assume you have the “right” .torrent file
- ◆ Can peers modify parts of the downloaded file(s)?
- ◆ Can you be tracked?

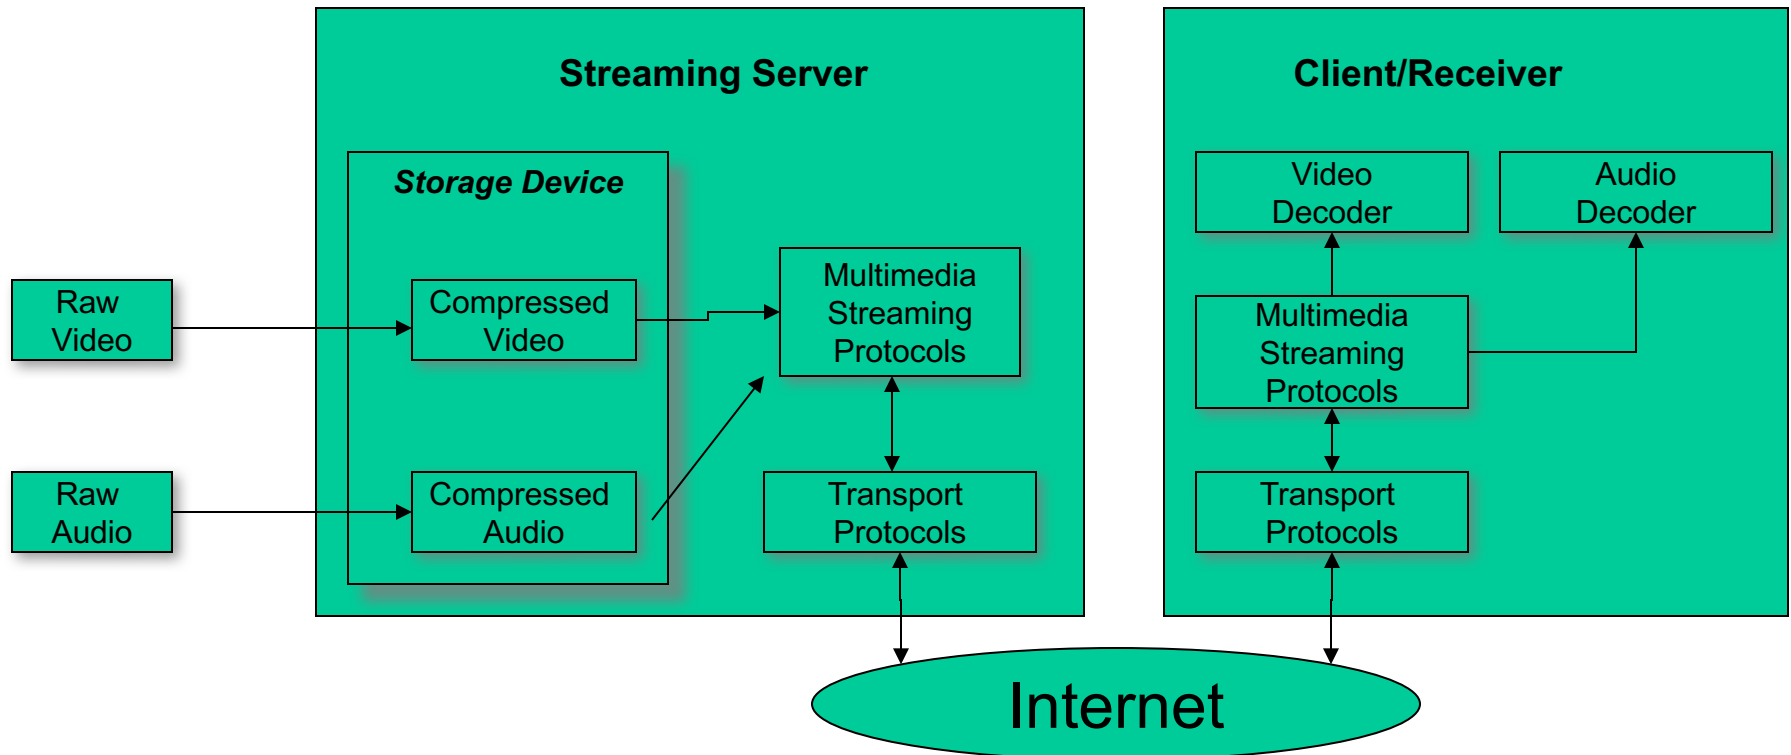
Internet Video

What is Multimedia Streaming?

- **Multimedia Streaming:** Clients request audio/video files from servers and pipeline reception over the network and display

User's perspective:

- Quick start without waiting for full download.
- Coming continuously without interruption.
- VCR operation (pause, resume, fast forward, rewind, etc.)



Challenges in Media Streaming Protocols

1. Rate Control:
Determine the sending rate
based on the available
bandwidth in the network.

Clients/Receivers



Streaming Server



Ethernet

Broadband/LTE



2G

3. Continuous
Distribution: TCP/UDP/IP
suite provides best-effort,
no guarantees on
expectation or variance of
packet delay

2. Error Control:
Improve video
presentation quality in the
presence of packet loss.



Techniques in Multimedia Streaming Protocols (1)

- ◆ Rate Control
 - Scalable compression
 - Base substream and enhancement substreams.
 - SNR scalability / spatial scalability / temporal scalability
 - Rate filter
 - Frequency filter
 - Frame-dropping filter
 - Re-quantization filter
 - QoS Feedback, e.g. RTCP.



Techniques in Multimedia Streaming Protocols (2)

- Error Control
 - Add redundant data in coding
 - MDC, (Multiple Description Coding)
 - FEC (Forward Error Coding)
 - Receiver End Error Concealment
 - Receiver conceal data loss.
 - Spatial interpolation, used in intra-coded frame.
 - Temporal interpolation, used in inter-coded frame.

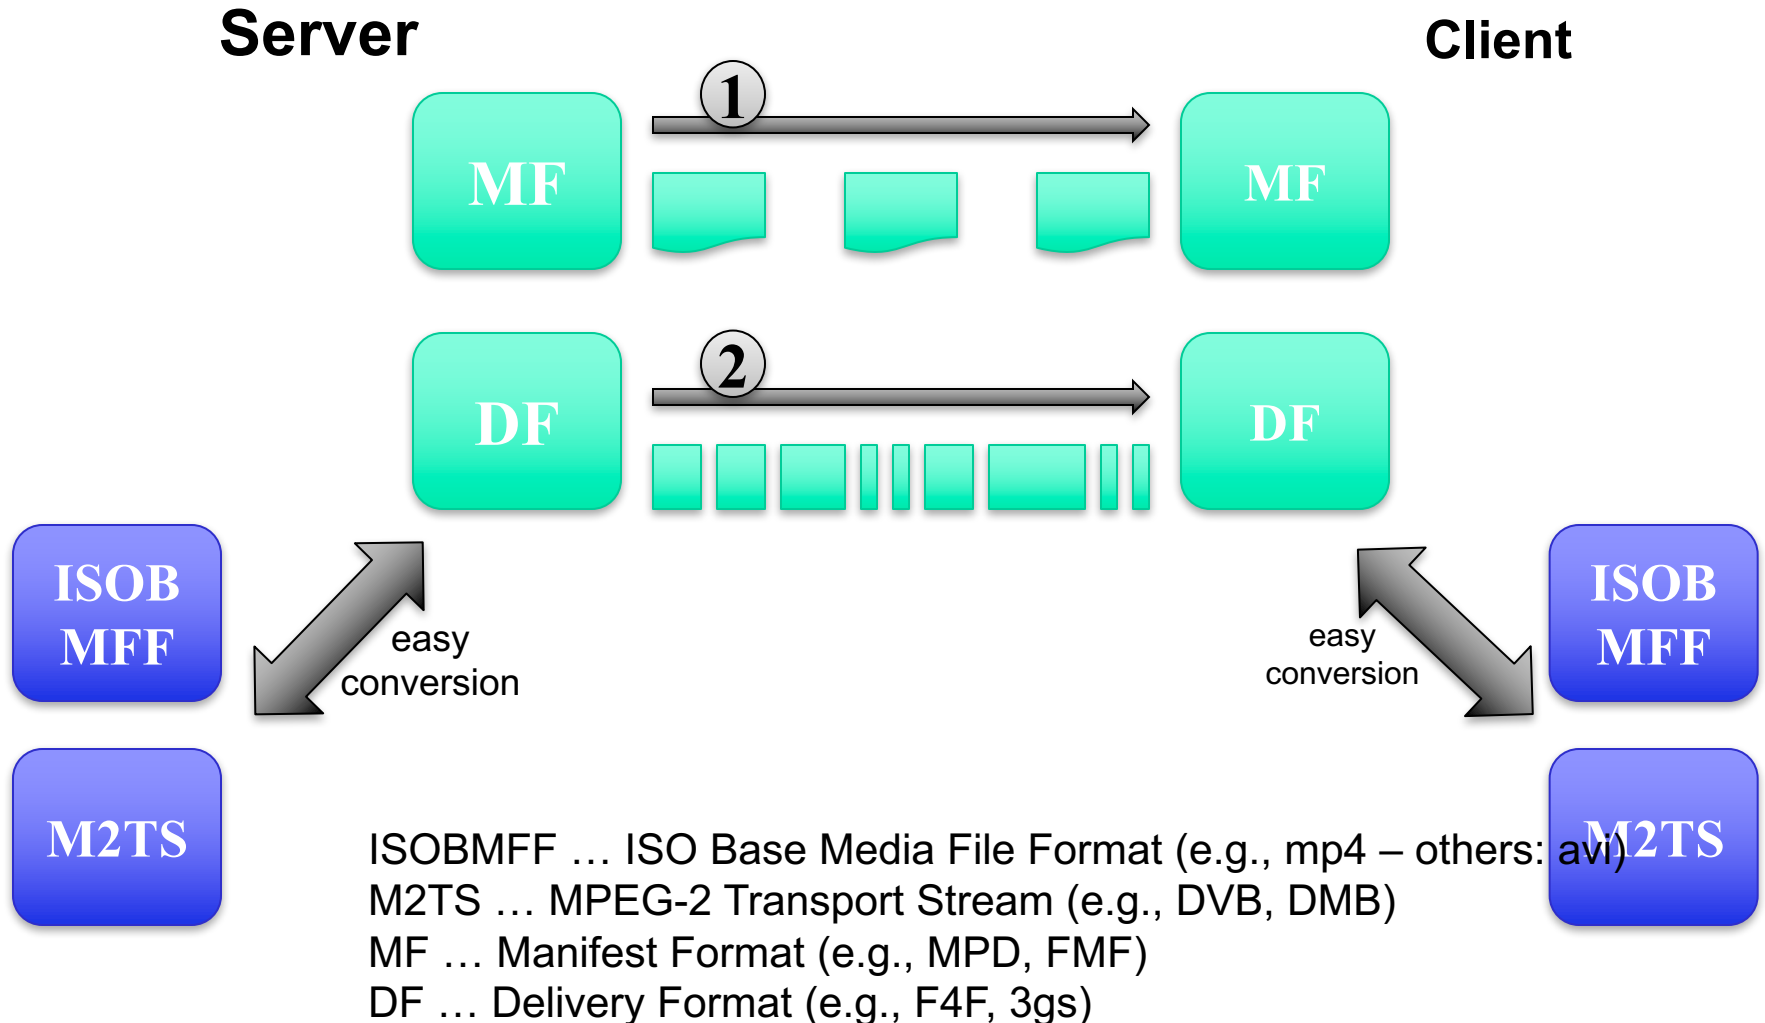
A General View of the Classic Internet Multimedia Streaming Protocols

- ◆ **Stream description** SDP, SMIL...
Describe the session and content
- ◆ **Stream control** RTSP
Remote control the session
- ◆ **Media transport** RTP
Error control and flow control
- ◆ **Resource reservation** (if any!): RSVP, DiffServ
provide QoS for media streaming packets

HTTP-Based Internet Video

Use HTTP to scale video
distribution

HTTP Streaming of Media



Adaptive Streaming in Practice



HTTP Live Streaming



Microsoft
Silverlight™

Smooth Streaming



Adobe HTTP Dynamic
Streaming

Combined A/V
streams only



MPEG DASH


Separate Audio/Video

Ack & ©: Mark Watson
2010/05/02

MPEG DASH Data Model

Media Presentation Description (MPD) Data Model

- ▶ MPD describes accessible Segments and corresponding timing



Media Presentation Description

- ◆ Redundant information of Media Streams for the purpose to initially select or reject Groups or Representations
 - Examples: Codec, DRM, language, resolution, bandwidth
- ◆ Access and Timing Information
 - HTTP-URL(s) and byte range for each accessible Segment
 - Earliest next update of the MPD on the server
 - Segment availability start and end time in wall-clock time
 - Approximated media start time and duration of a Media Segment in the media presentation timeline
 - For live service, instructions on starting playout such that media segments will be available in time for smooth playout in the future
- ◆ Switching and splicing relationships across Representations
- ◆ Relatively little other information