

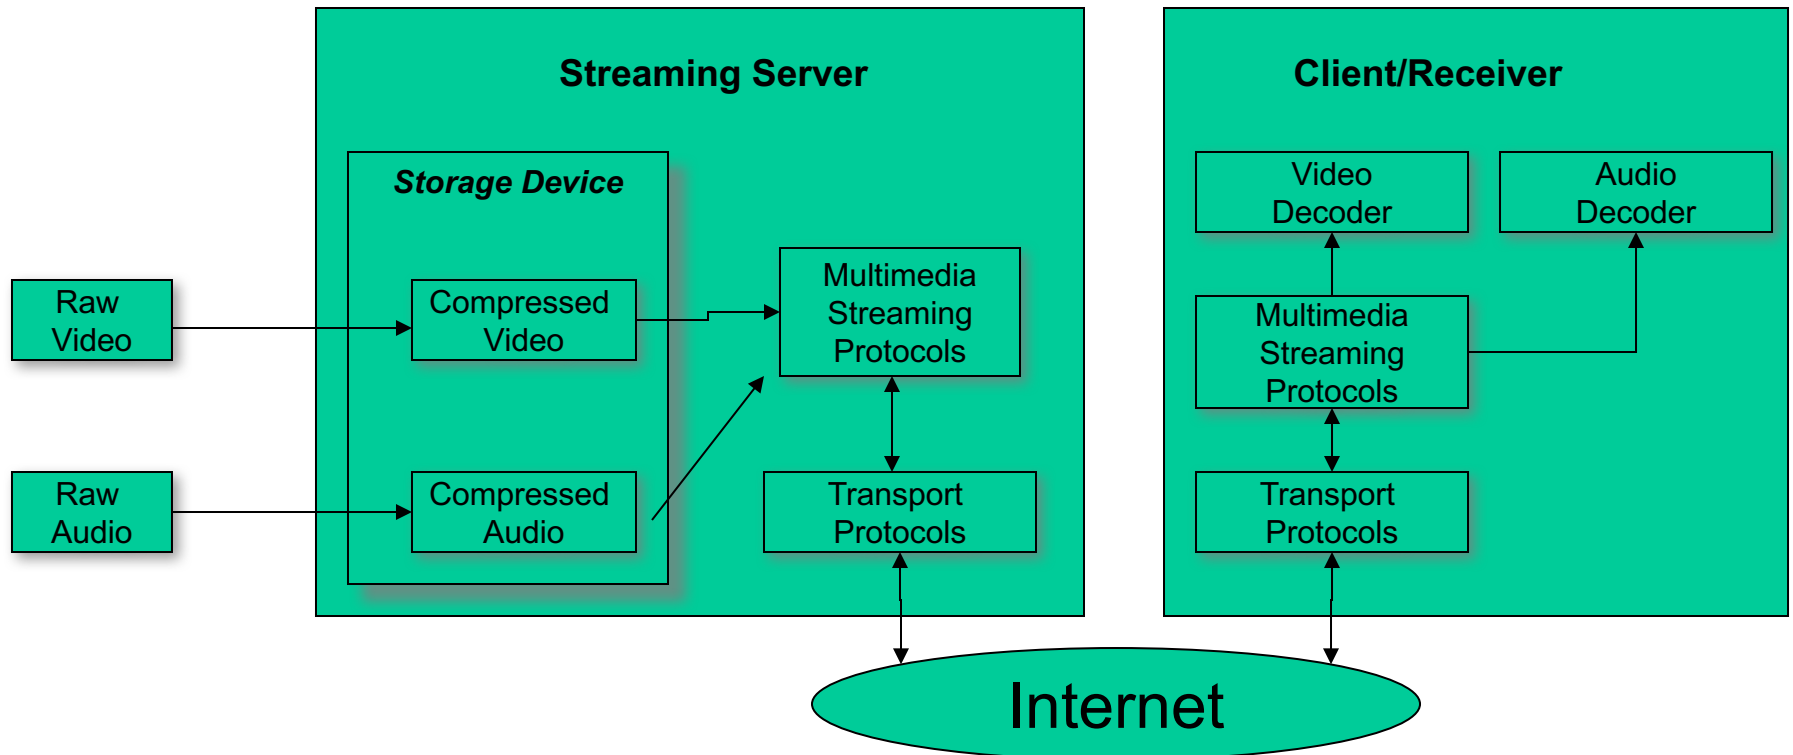
Internet Video

What is Multimedia Streaming?

- **Multimedia Streaming:**
Clients request audio/video files from servers and pipeline reception over the network and display

User's perspective:

- Quick start without waiting for full download.
- Coming continuously without interruption.
- VCR operation (pause, resume, fast forward, rewind, etc.)



Challenges in Media Streaming Protocols

1. Rate Control:
Determine the sending rate
based on the available
bandwidth in the network.

Clients/Receivers



Streaming Server



Ethernet

Broadband/LTE



2G

3. Continuous
Distribution: TCP/UDP/IP
suite provides best-effort,
no guarantees on
expectation or variance of
packet delay

2. Error Control:
Improve video
presentation quality in the
presence of packet loss.



Techniques in Multimedia Streaming Protocols (1)

- ◆ Rate Control
 - Scalable compression
 - Base substream and enhancement substreams.
 - SNR scalability / spatial scalability / temporal scalability
 - Rate filter
 - Frequency filter
 - Frame-dropping filter
 - Re-quantization filter
 - QoS Feedback, e.g. RTCP.



Techniques in Multimedia Streaming Protocols (2)

- Error Control
 - Add redundant data in coding
 - MDC, (Multiple Description Coding)
 - FEC (Forward Error Coding)
 - Receiver End Error Concealment
 - Receiver conceal data loss.
 - Spatial interpolation, used in intra-coded frame.
 - Temporal interpolation, used in inter-coded frame.

A General View of the Classic Internet Multimedia Streaming Protocols

- ◆ Stream description SDP, SMIL...

Describe the session and content

- ◆ Stream control RTSP

Remote control the session

- ◆ Media transport RTP

Error control and flow control

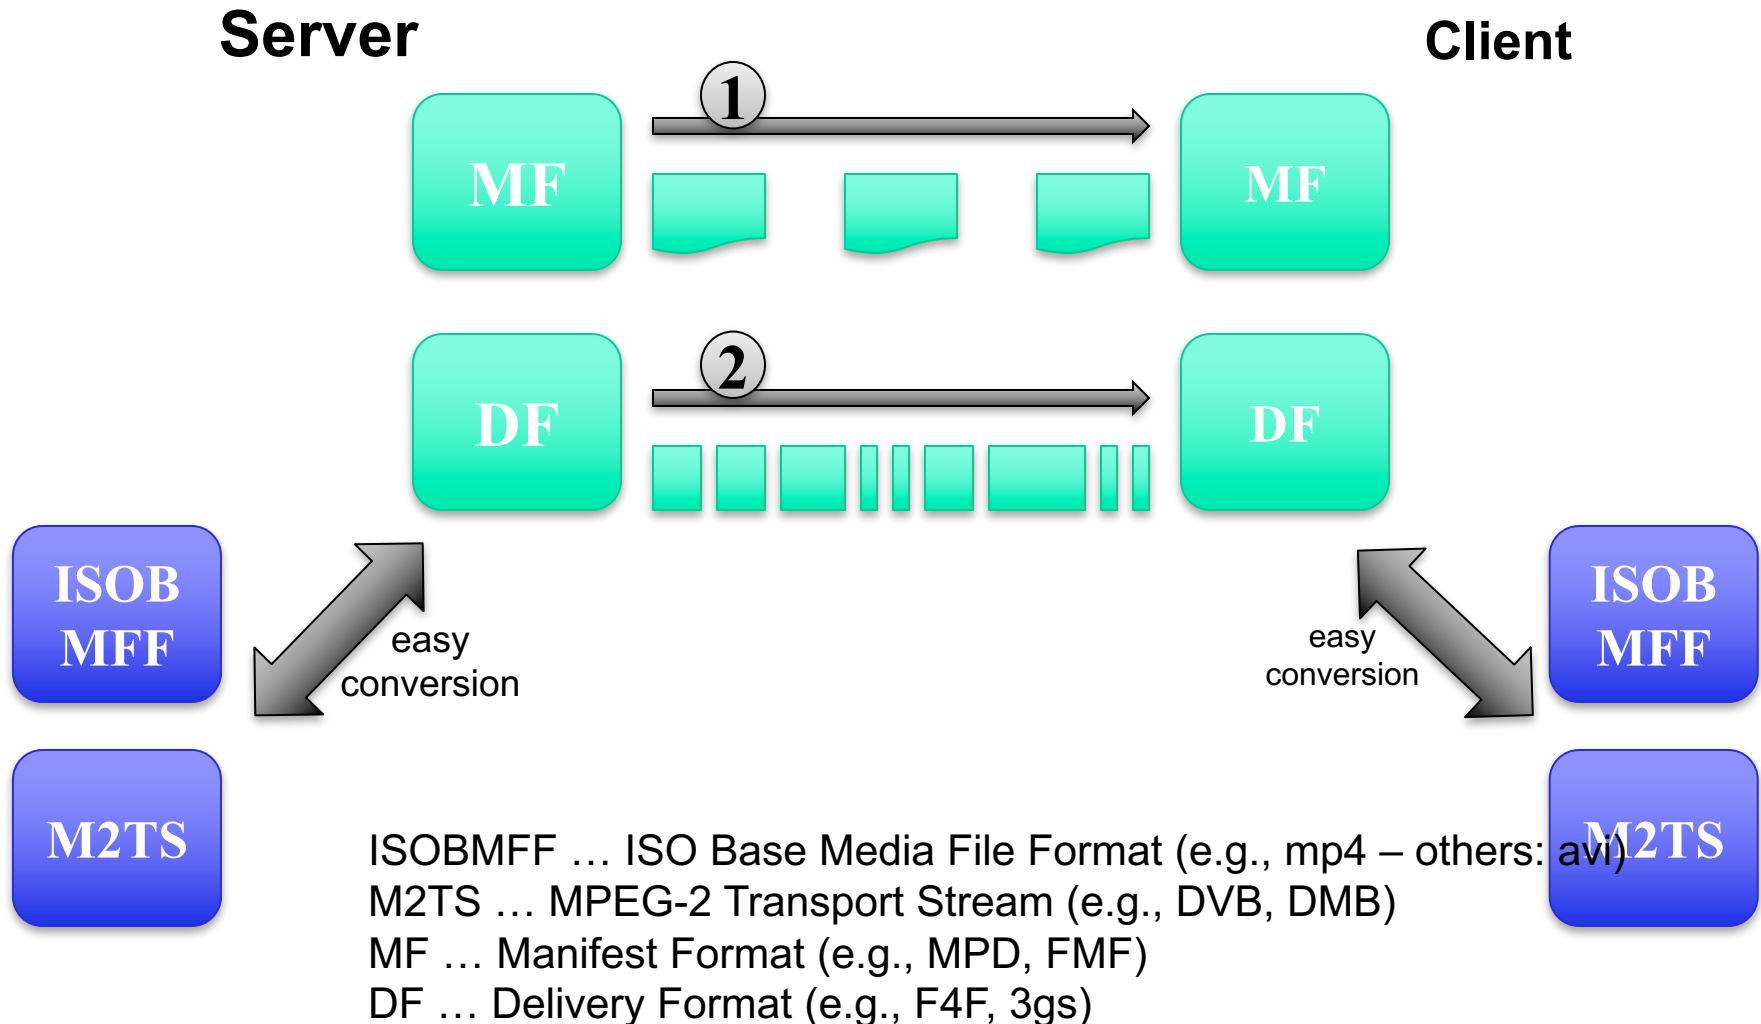
- ◆ Resource reservation (if any!): RSVP, DiffServ

provide QoS for media streaming packets

HTTP-Based Internet Video

Use HTTP to scale video
distribution

HTTP Streaming of Media



Adaptive Streaming in Practice



HTTP Live Streaming



Microsoft
Silverlight™

Smooth Streaming



Adobe HTTP Dynamic
Streaming

Combined A/V
streams only



MPEG DASH

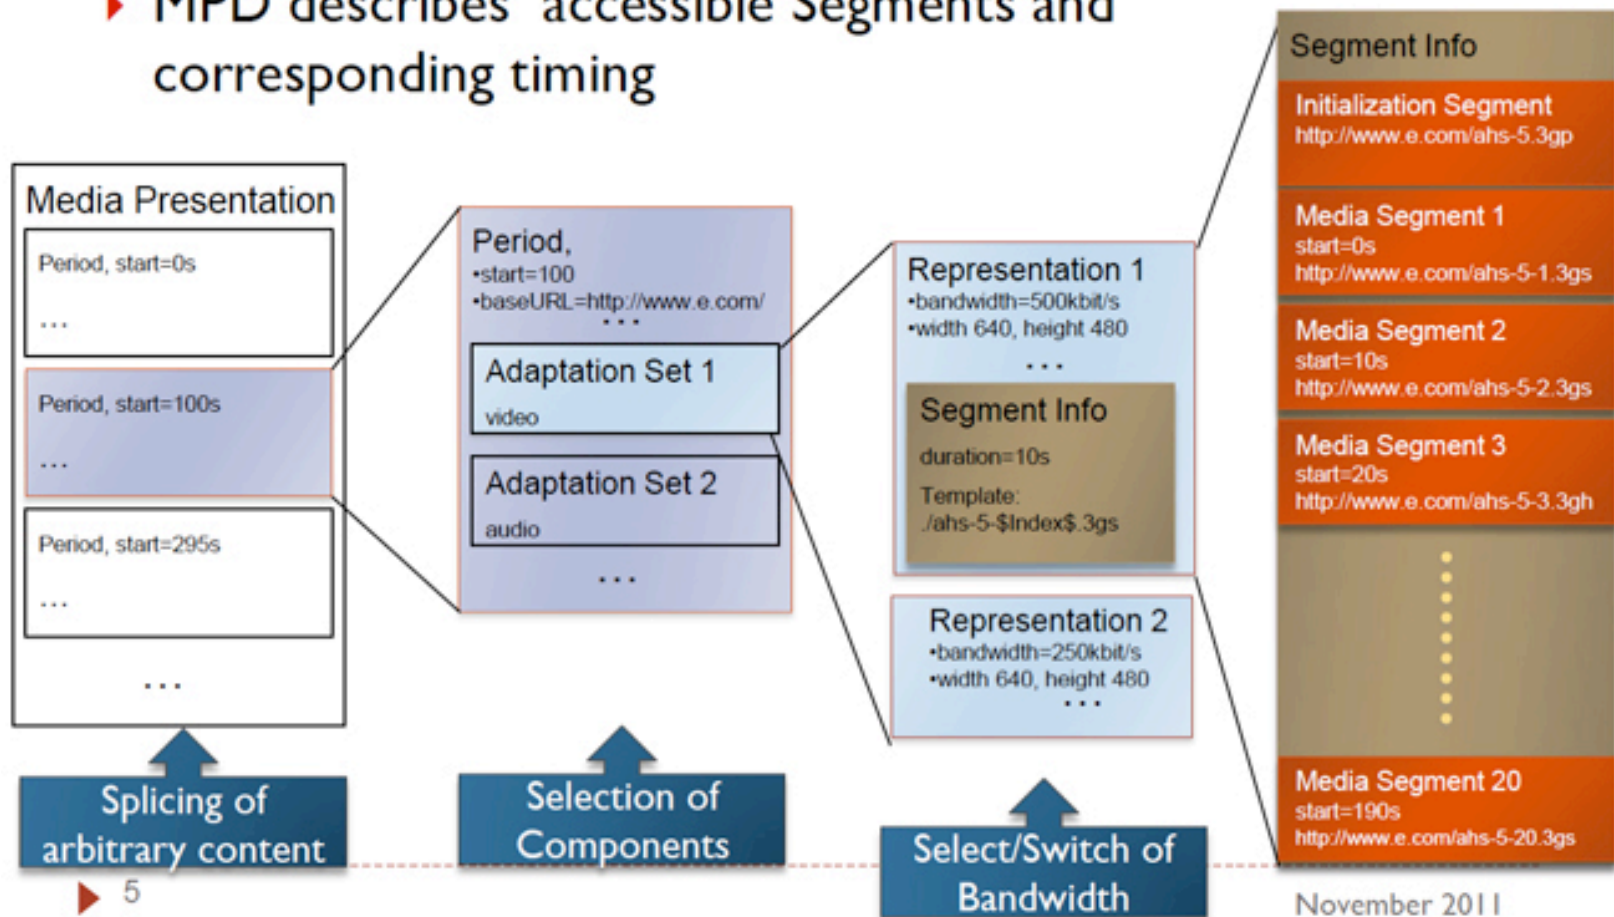
Separate Audio/Video

Ack & ©: Mark Watson
2010/05/02

MPEG DASH Data Model

Media Presentation Description (MPD) Data Model


- ▶ MPD describes accessible Segments and corresponding timing



Media Presentation Description

- ◆ **Redundant** information of **Media Streams** for the purpose to initially select or reject Groups or Representations
 - Examples: Codec, DRM, language, resolution, bandwidth
- ◆ **Access and Timing Information**
 - **HTTP-URL(s)** and **byte range** for each accessible Segment
 - Earliest next update of the MPD on the server
 - Segment **availability start and end time** in wall-clock time
 - Approximated **media start time and duration** of a Media Segment in the media presentation timeline
 - For **live service**, instructions on starting playout such that media segments will be available in time **for smooth playout** in the future
- ◆ Switching and splicing relationships across Representations
- ◆ Relatively little other information

Moving to Transport Layer



Chapter 3

3.1 Transport-layer services

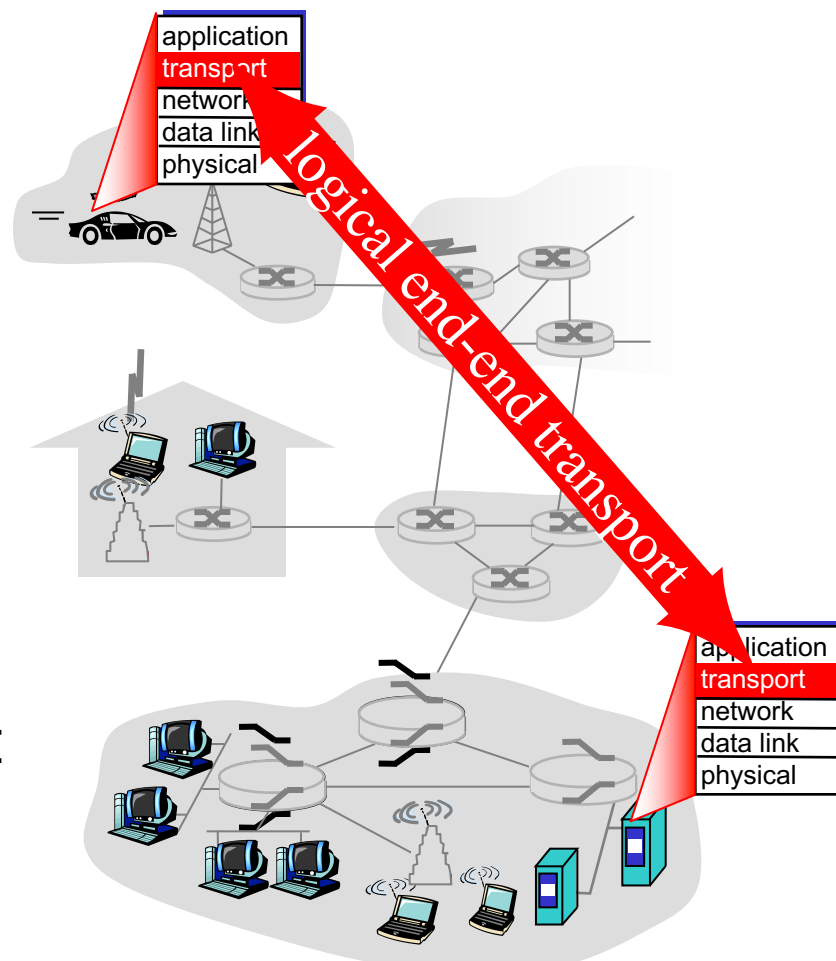
3.2 Multiplexing and de-multiplexing

3.3 Connectionless transport: UDP

3.4 Principles of reliable data transfer

Transport Layer

- ◆ Provide *logical communication* between app processes
- ◆ transport protocols run in end systems
 - send side: breaks app data into **segments**, passes to network layer
 - rcv side: reassembles segments into original data format, passes to app layer
- ◆ Multiple transport protocols exist
 - Mostly used: TCP, UDP
 - Other transports: RTP, SCTP
 - Recent ones: QUIC, SPDY, HTTP/2



Check your
/etc/protocols

Transport vs. network layer

- ♦ *transport layer*: logical communication between processes
 - relies on, enhances, network layer services
- ♦ *network layer*: logical communication between hosts

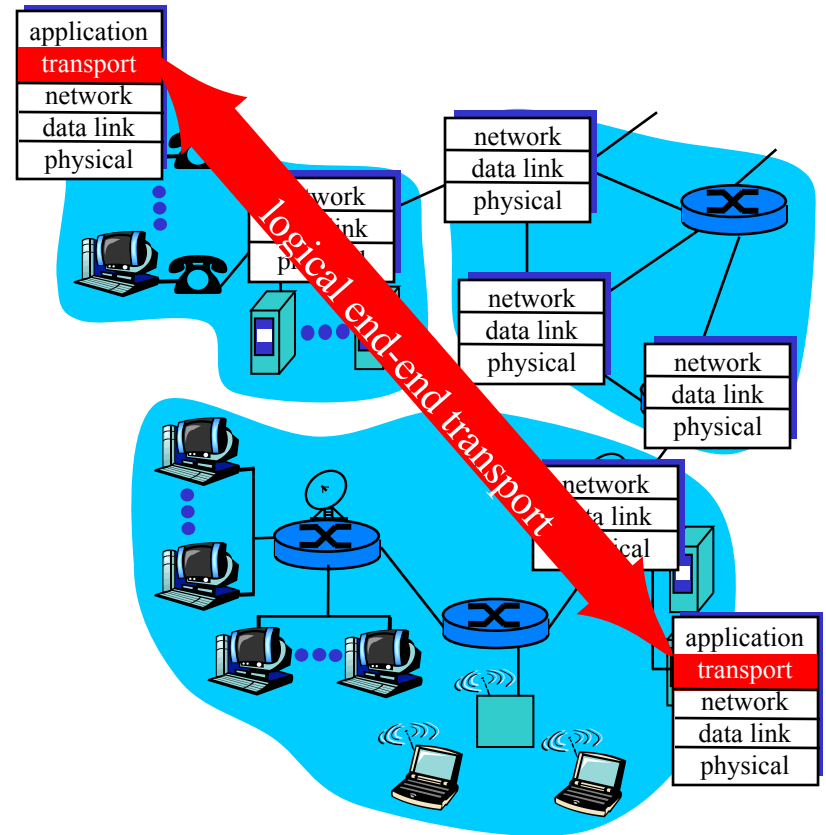
Household analogy:

12 young kids sending letters to 12 other kids

- ♦ processes = kids
- ♦ app messages = letters in envelopes
- ♦ hosts = houses
- ♦ transport protocol = kids parents
- ♦ network-layer protocol = postal service

Internet transport-layer protocols

- ◆ Reliable, in-order delivery (TCP)
 - congestion control
 - flow control
 - connection setup
- ◆ Unreliable, unordered delivery: UDP
- ◆ FYI:
 - RTP: unreliable delivery with timing information
 - SCTP: reliable delivery of multiple substreams



Common function among all the above:
multiplexing/demultiplexing

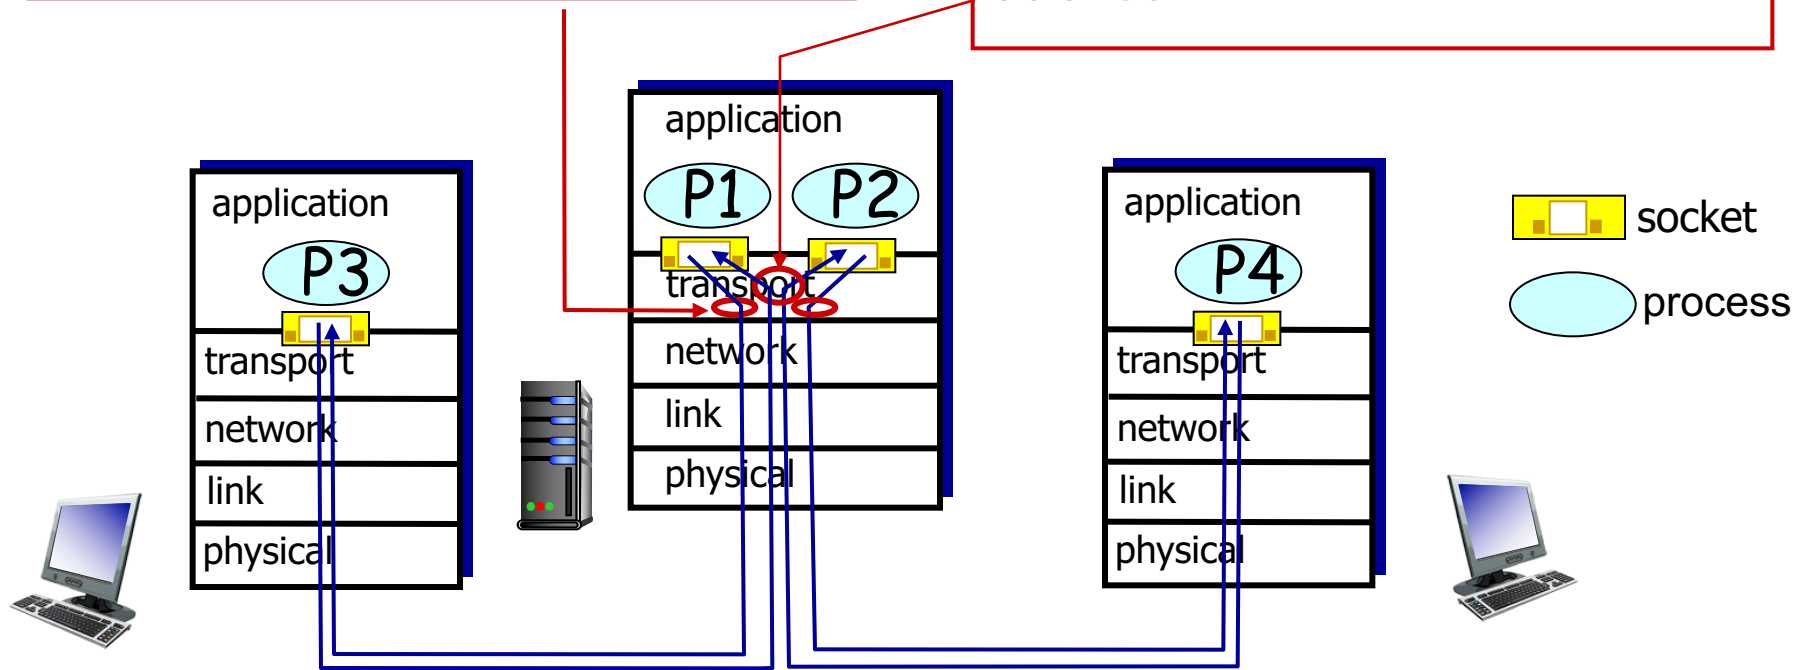
Multiplexing/demultiplexing

multiplexing at sender:

handle data from multiple sockets, add transport header (later used for demultiplexing)

demultiplexing at receiver:

use header info to deliver received segments to correct socket




Each process is identified by IP address and port#

How Demultiplexing Works

◆ Host receives IP datagrams

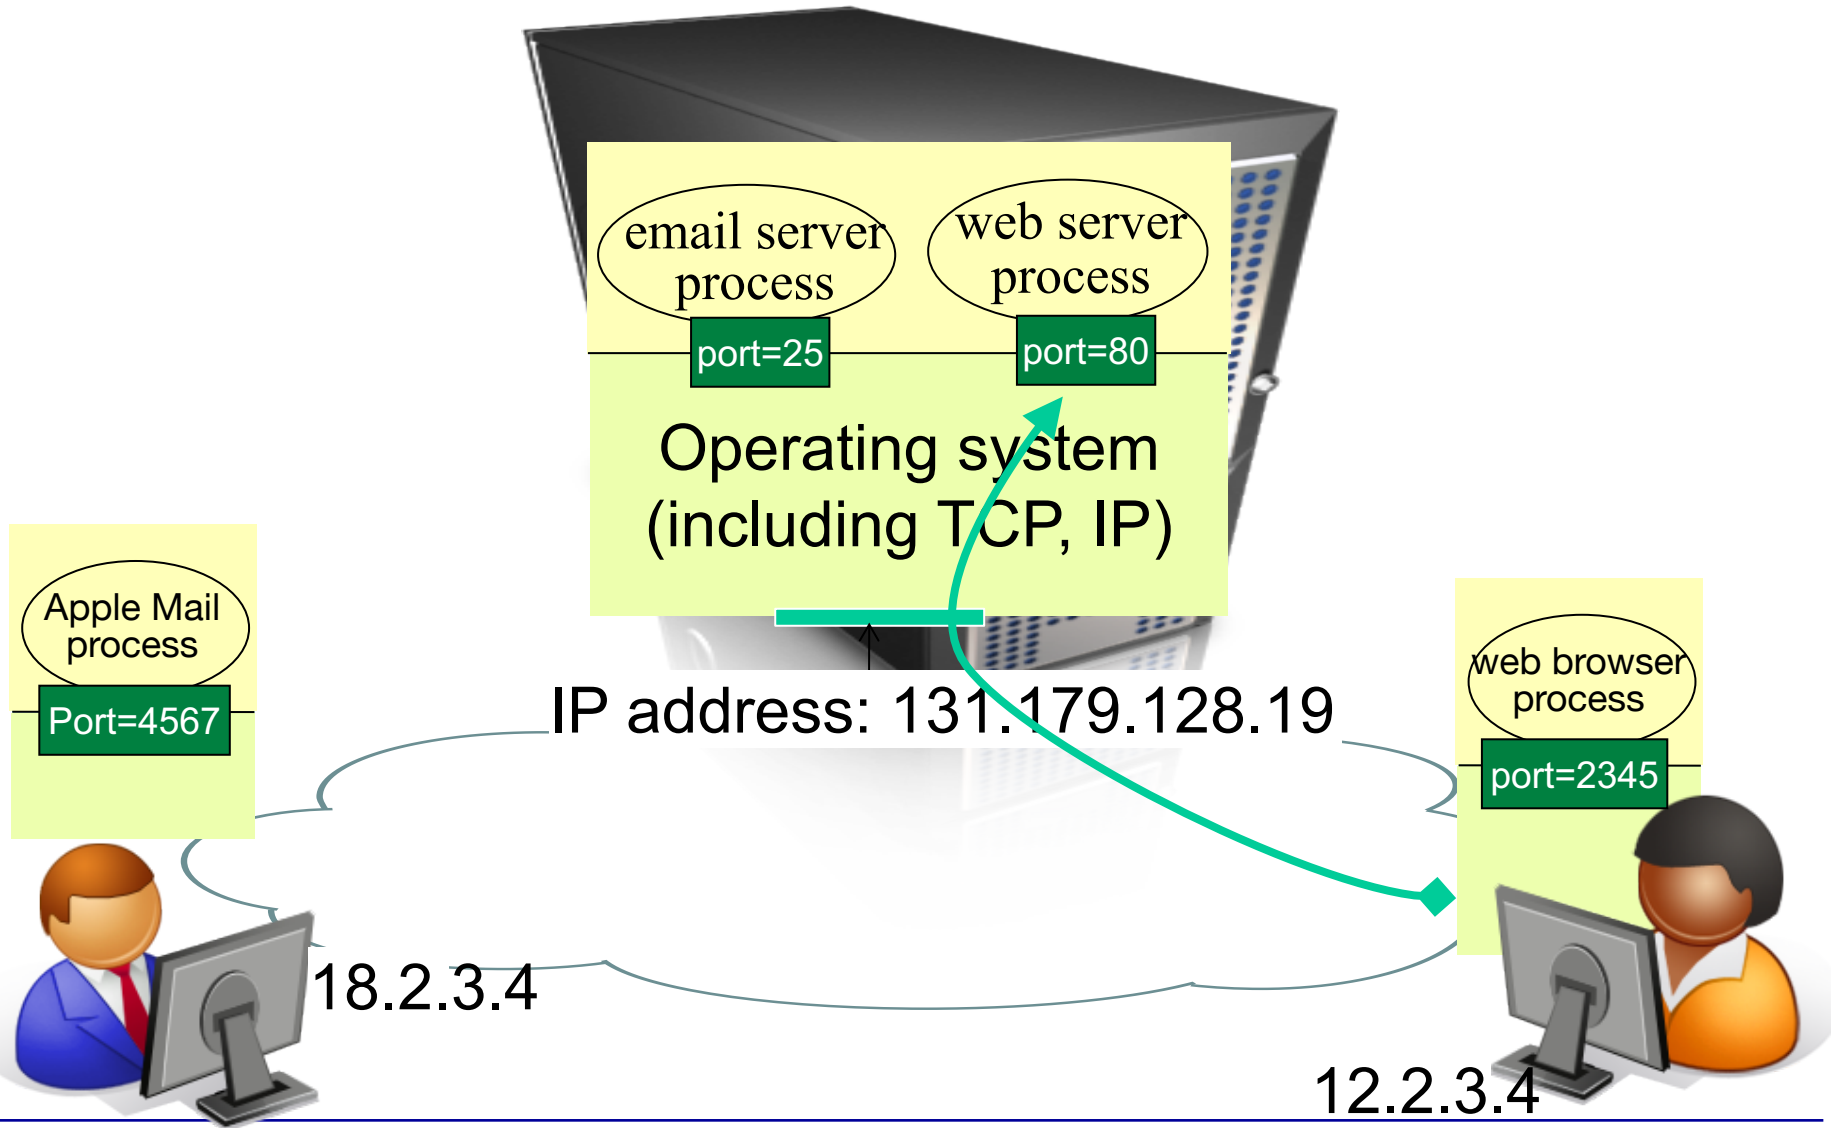
- each datagram has source IP address, destination IP address
- each datagram carries one transport-layer segment
- each segment has source, destination port number

◆ Host uses IP addresses & port numbers to direct segment to appropriate socket



TCP/UDP segment format

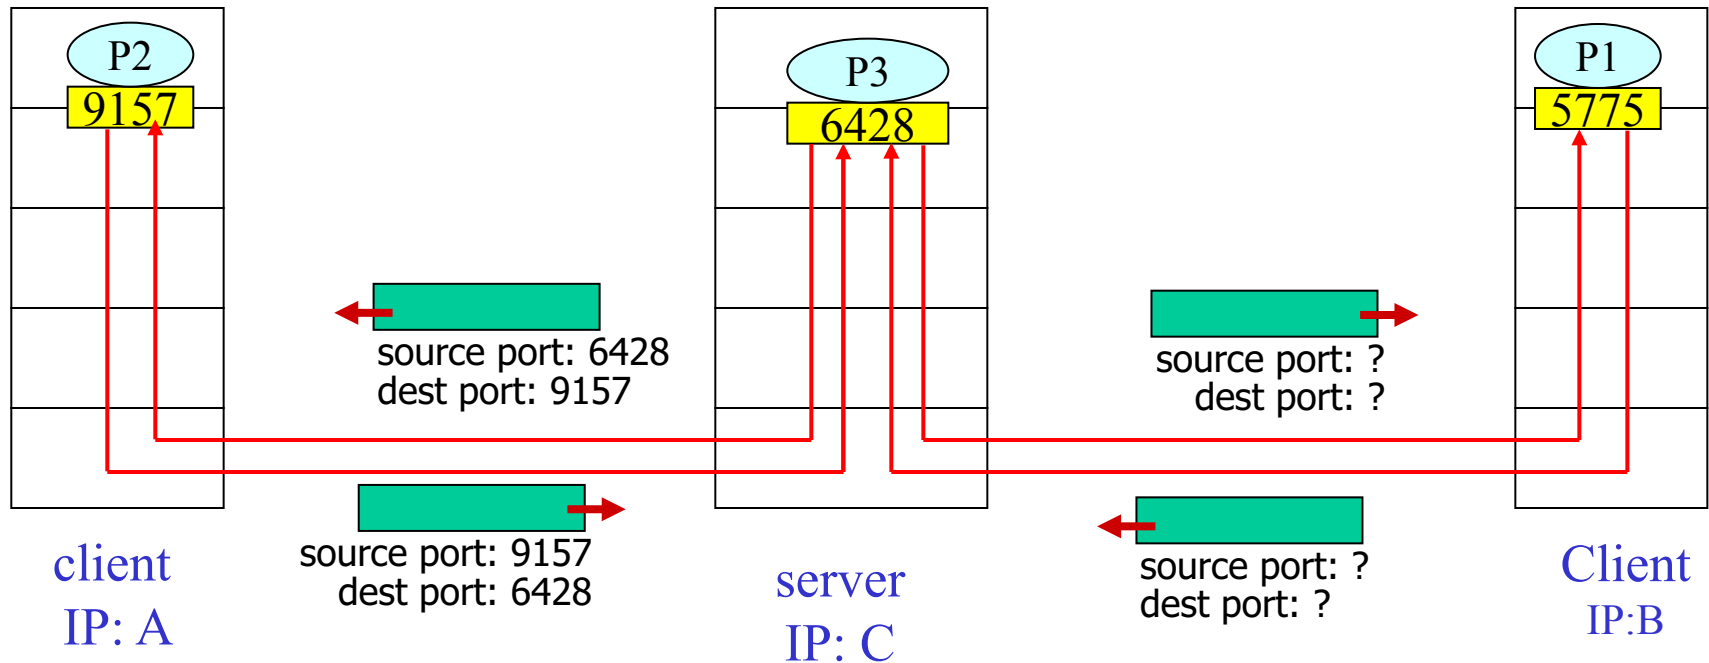
IP address, TCP connection, port number, processes, and sockets



Connectionless Demultiplexing

- ◆ when creating a packet to send to a UDP socket, must specify
 - destination IP address
 - destination port #
- ◆ When host receives a UDP packet:
 - checks destination port# in the packet
 - directs the packet to socket with that port#
- ◆ IP packets with *same dest. port #* are directed to the same socket at the destination host
 - They may have different source IP addresses and/or source port#s

Connectionless demux (cont)



How can a server know where to return a reply?

FYI: From UDP spec (RFC768) “UDP module must be able to determine the source and destination internet addresses and the protocol field from the internet header. One possible UDP/IP interface would return the whole internet datagram including all of the internet header in response to a receive operation”

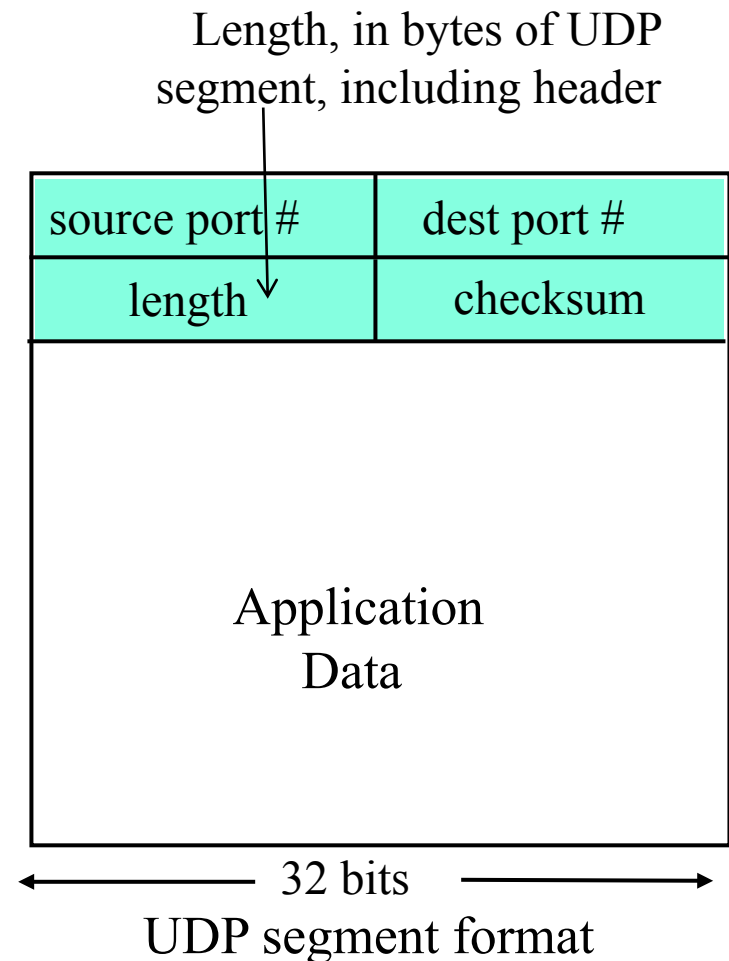
UDP: User Datagram Protocol [RFC 768]

- ◆ Best effort service: UDP segments may be lost, duplicated, or delivered out of order to app. processes
- ◆ *connectionless*:
 - no prior handshaking between UDP sender, receiver
 - each UDP segment handled independently of others
- ◆ UDP use:
 - DNS
 - streaming multimedia apps (loss tolerant, rate sensitive)
- ◆ For reliable transfer: add reliability at application layer

UDP header format

Why is there a UDP?

- ◆ no connection establishment
- ◆ simple: no connection state at sender, receiver
- ◆ small header size
- ↓ no congestion control: UDP can blast away as fast as desired



UDP checksum

Goal: detect bit errors in transmitted segment

Sender:

- ◆ treat segment contents as sequence of 16-bit integers
- ◆ checksum: addition (1's complement sum) of segment contents
- ◆ sender puts checksum value into UDP checksum field

Receiver:

- ◆ compute checksum of received segment
- ◆ check if computed checksum equals checksum field value:
 - NO - error detected
 - YES - no error detected

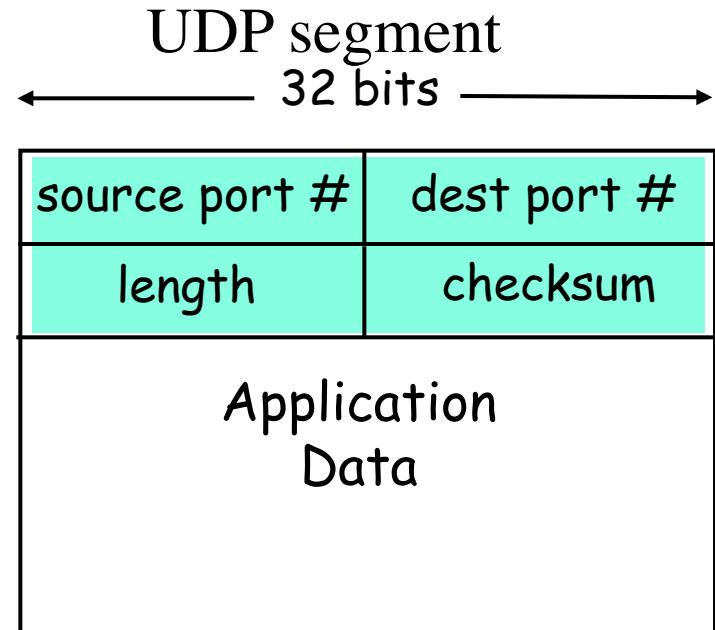
UDP Checksum

- ◆ https://en.wikipedia.org/wiki/User_Datagram_Protocol#Checksum_computation
- ◆ The 16-bit checksum field is used for error-checking received packet
- ◆ **16-bit one's complement of the one's complement sum of a pseudo header of information from the IP header, the UDP header, and the data, padded with zero octets at the end (if necessary) to make a multiple of two octets**
 - ◆ *For computing, the value of the checksum field is zero*

How to Calculate UDP Checksum

◆ UDP header

- Length: # of bytes (including both header & data)
- checksum: computed over
 - the **pseudo header**, and
 - UDP header and data.
 - if checksum field=0: no checksum



- ◆ **pseudo header**: UDP's self-protection against misdelivered IP packets
pseudo header is not carried in UDP packet, nor counted in the length field

source IP address		
destination IP address		
zero	protocol	UDP length

IPv4 Checksum

- ◆ https://en.wikipedia.org/wiki/IPv4#Header_Checksum
- ◆ The 16-bit checksum field is used for error-checking of the IPv4 header
 - If value of the carried checksum don't match the calculated, the router discards the packet
 - Different from TCP and UDP checksums
- ◆ ***16-bit one's complement of the one's complement sum of all 16-bit words in the IPv4 ***header******
 - *For computing, the value of the checksum field is zero*

IPv4 Checksum Example

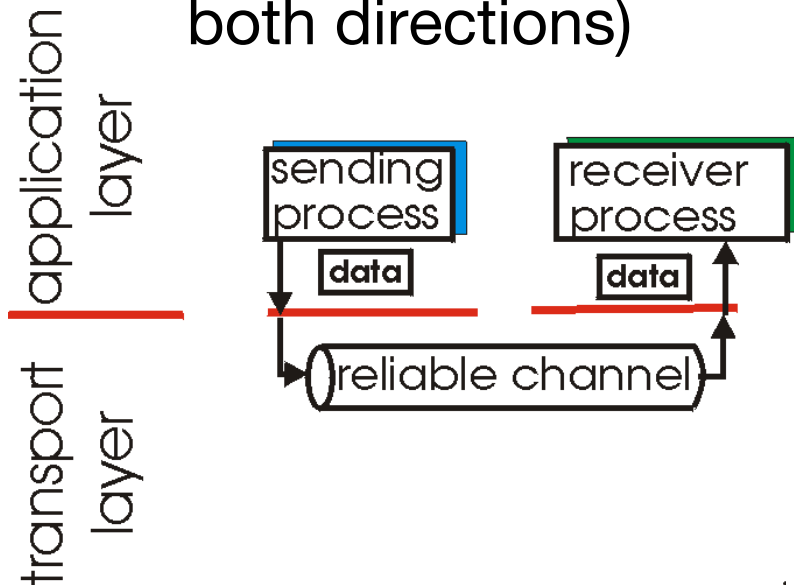
- ◆ 4500 0030 4422 4000 8006 0000 8C7C 19AC
AE24 1E2B₁₆ (20 bytes IP header)
- ◆ Sum up every 16 bits
 - $4500_{16} + 0030_{16} + 4422_{16} + 4000_{16} + 8006_{16} + 0000_{16} + 8C7C_{16} + 19AC_{16} + AE24_{16} + 1E2B_{16} = 0002BBCF$ (32-bit sum)
- ◆ 1's complement 16-bit sum (carry out)
 - $0002_{16} + BBCF_{16} = BBD1_{16} = 1011101111010001_2$
 - Can do this carry out while summing up
- ◆ 1's complement of 1's complement 16-bit sum
 - $\sim BBD1_{16} = 0100010000101110_2$

Chapter 3 outline

- ◆ 3.1 Transport-layer services
- ◆ 3.2 Multiplexing and demultiplexing
- ◆ 3.3 Connectionless transport: UDP
- ◆ 3.4 Principles of reliable data transfer
- ◆ 3.5 Connection-oriented transport: TCP
 - segment structure
 - reliable data transfer
 - flow control
 - connection management
- ◆ 3.6 Principles of congestion control
- ◆ 3.7 TCP congestion control

Principles of Reliable Data Transfer

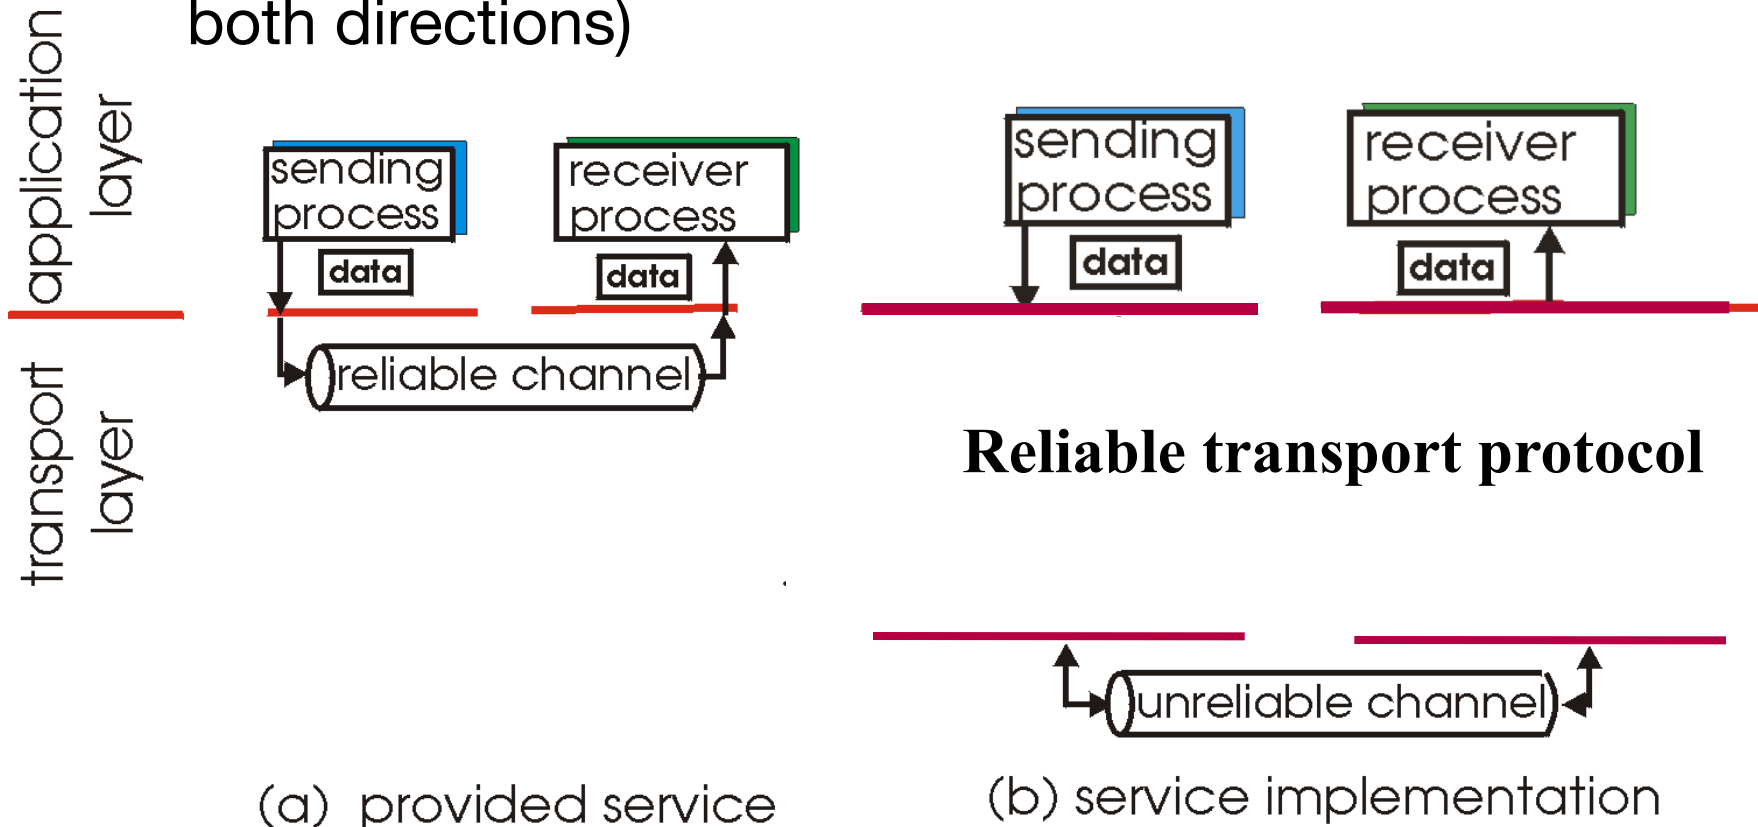
- ◆ characteristics of unreliable channel determines complexity of a reliable data transfer protocol (rdt)
- ◆ incrementally develop sender, receiver sides of rdt
 - consider one-way data transfer (control info will flow in both directions)



(a) provided service

Principles of Reliable Data Transfer

- ◆ characteristics of unreliable channel determines complexity of a reliable data transfer protocol (rdt)
- ◆ incrementally develop sender, receiver sides of rdt
 - consider one-way data transfer (control info will flow in both directions)



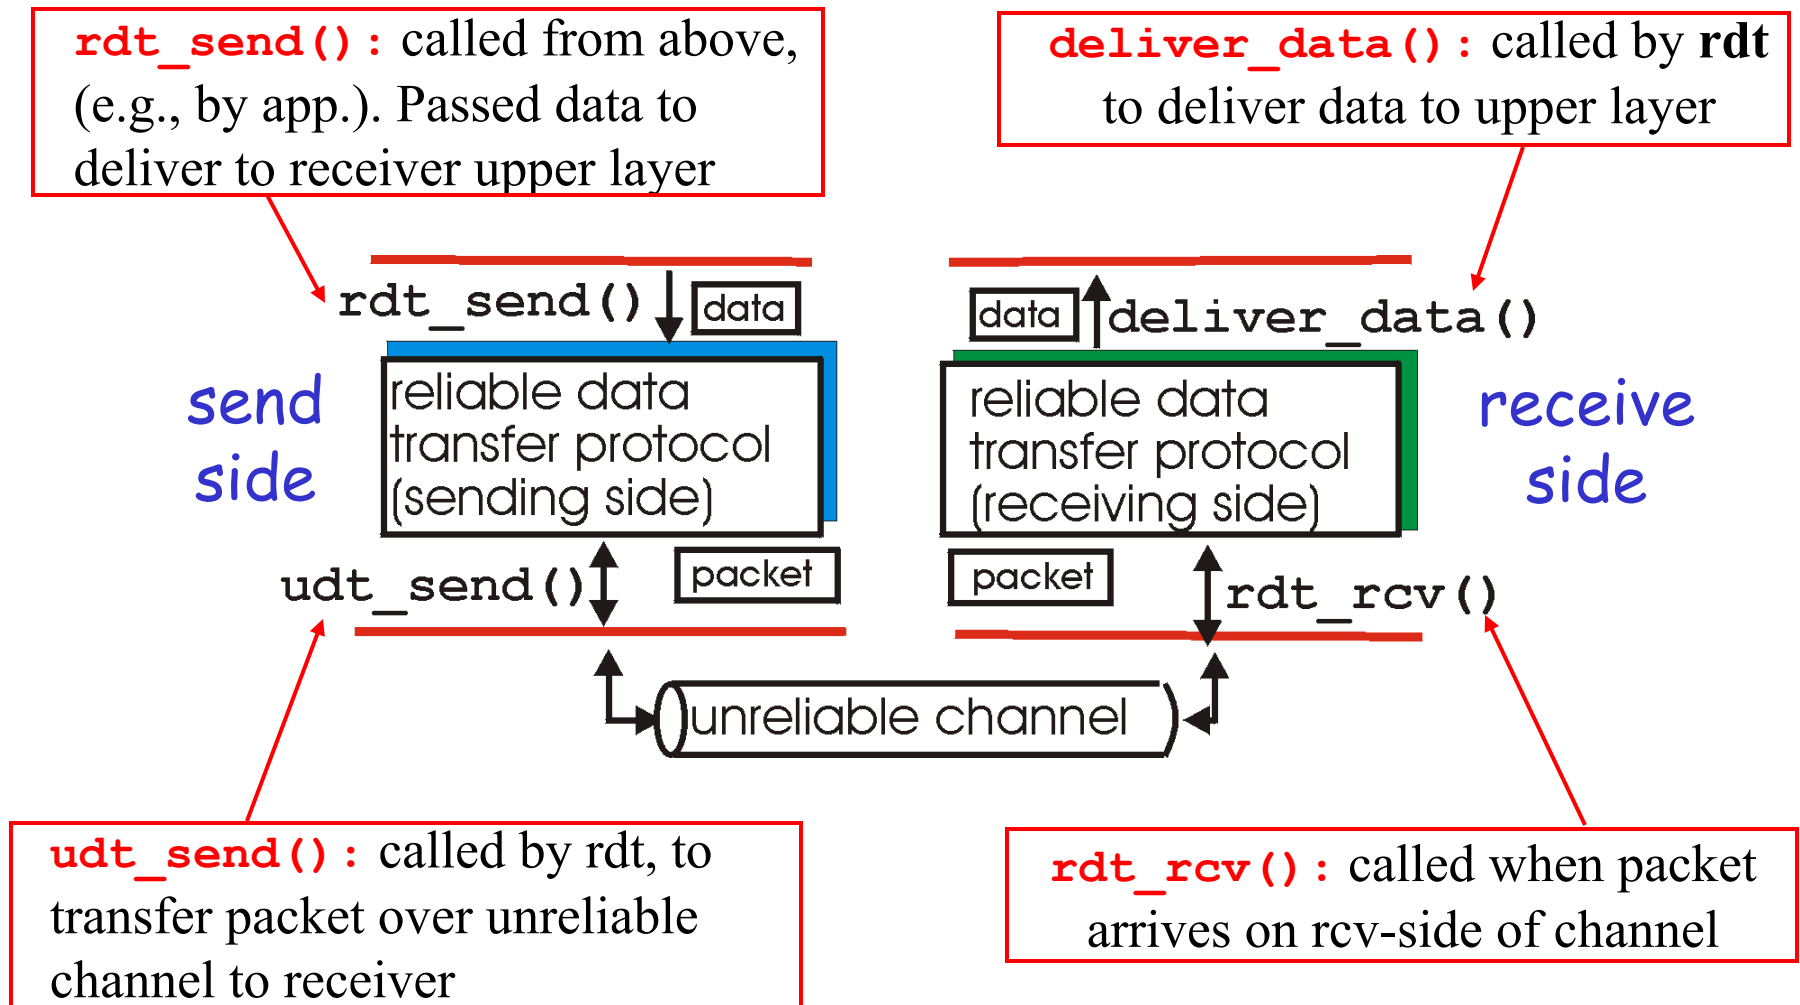
A simplified version of Principles of Reliable Data Transfer

- ◆ 3 questions
 - How many different types of errors?
 - How to detect each?
 - How to recover?
- ◆ 3 types of errors, and how to detect each
 - ◆ Corrupted bits in a packet: detected by checksum
 - ◆ Packet loss: detected by alarm timer (at sender end) in absence of ACK
 - ◆ Packets arrived out of order: detected by assigning each packet a sequence number
- ◆ Recovery: retransmitting the error/lost packet

Three basic components in reliable data delivery by retransmission

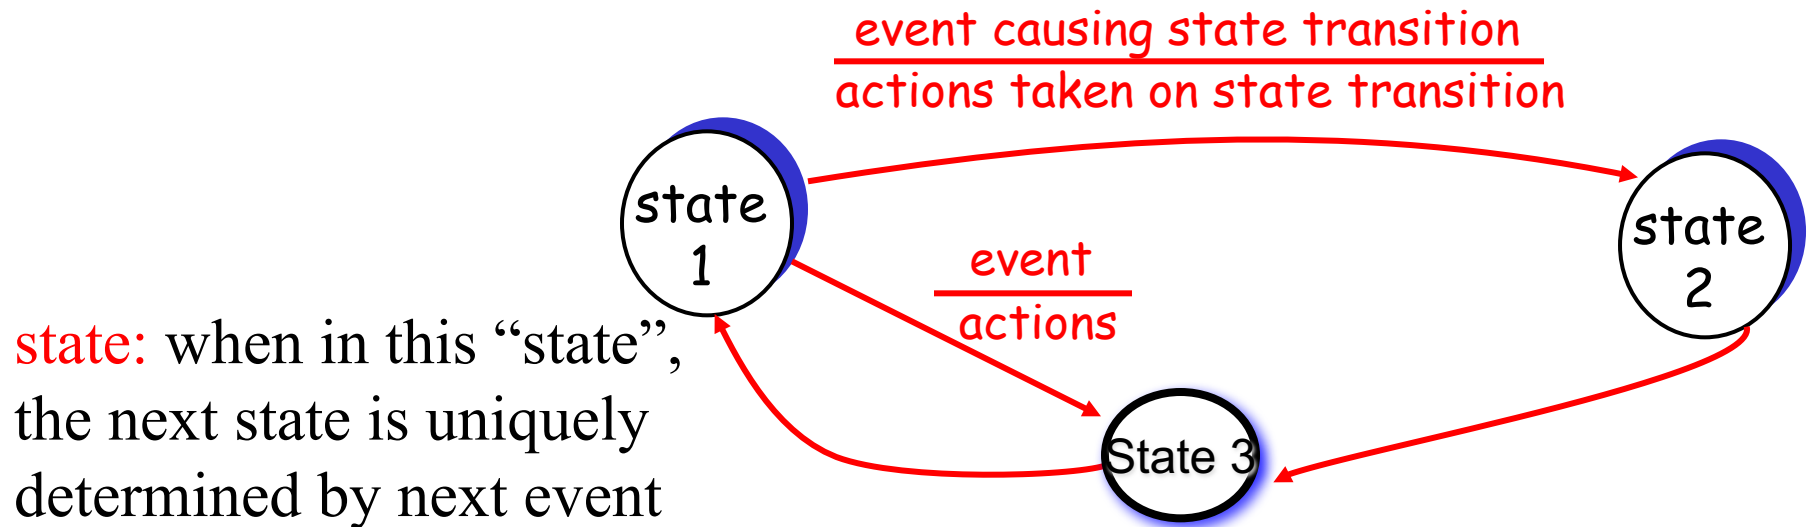
- ◆ **sequence #**: used to uniquely identify individual piece of data
- ◆ **Acknowledgment (ACK)**: reception report sent by receiver to the sender
- ◆ **Retransmission timer** set by the sender for the already sent, but has not been acknowledged packet
 - Retransmit the packet when timer expires

Reliable Data Transfer (rdt): getting started

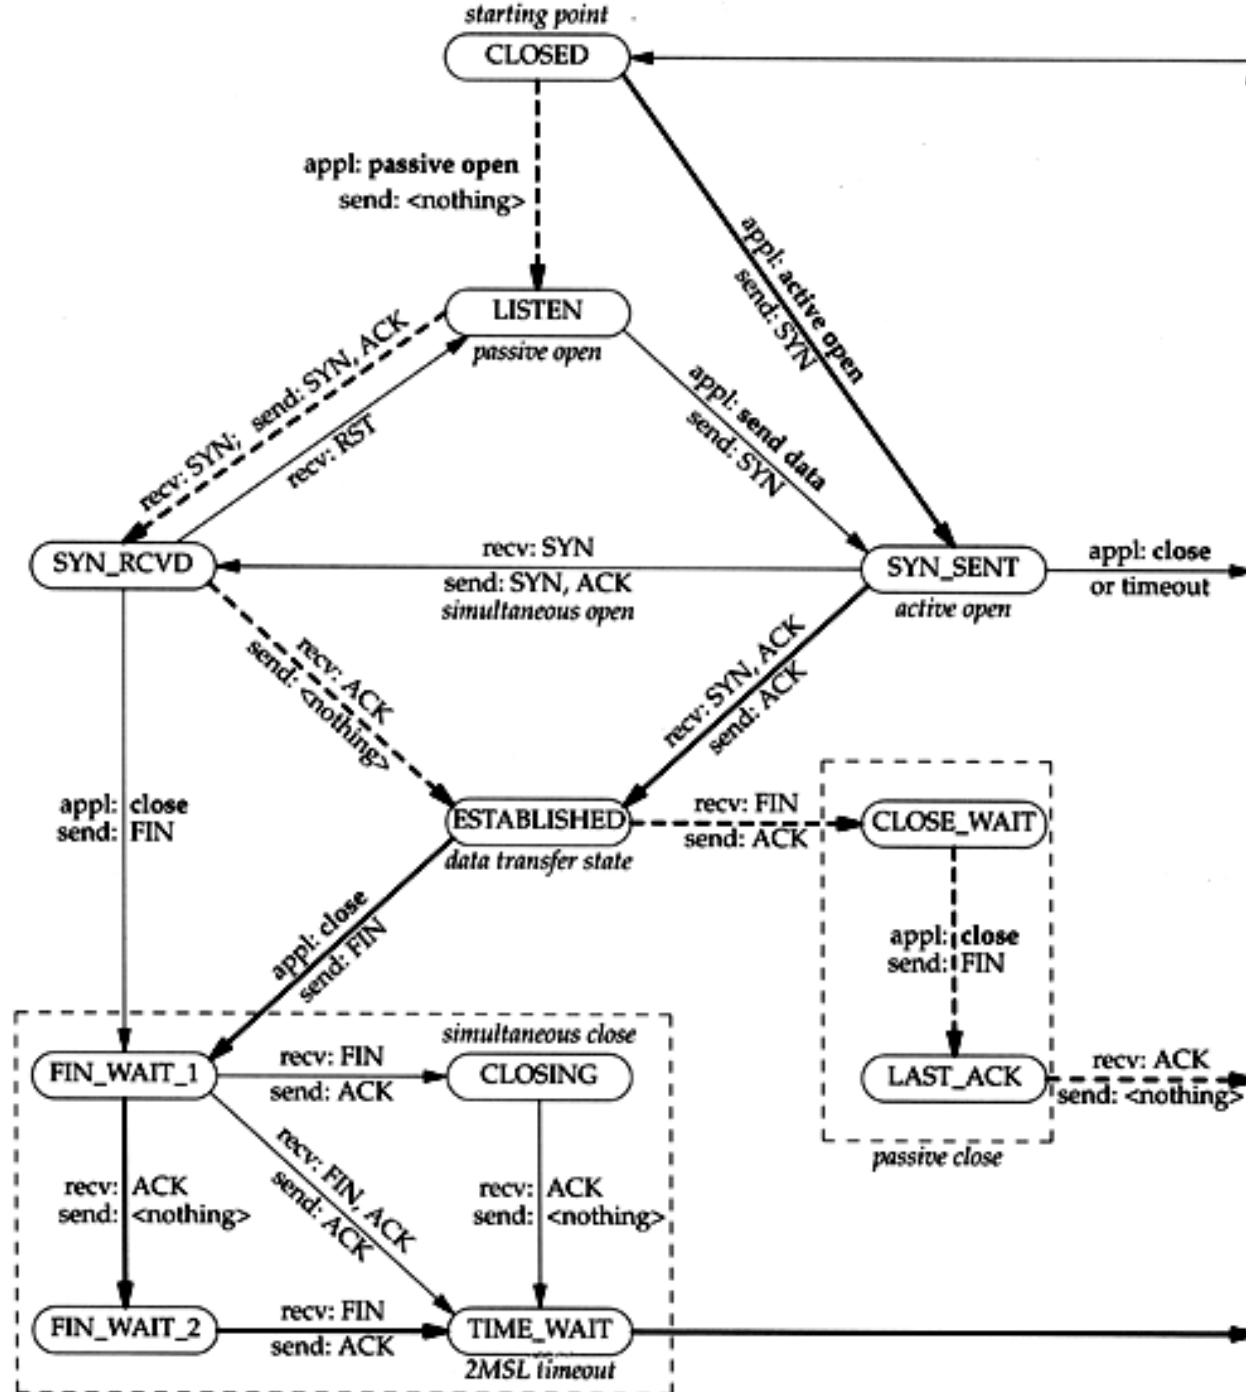


Reliable data transfer: getting started

- ◆ use finite state machines (FSM) to specify sender, receiver actions

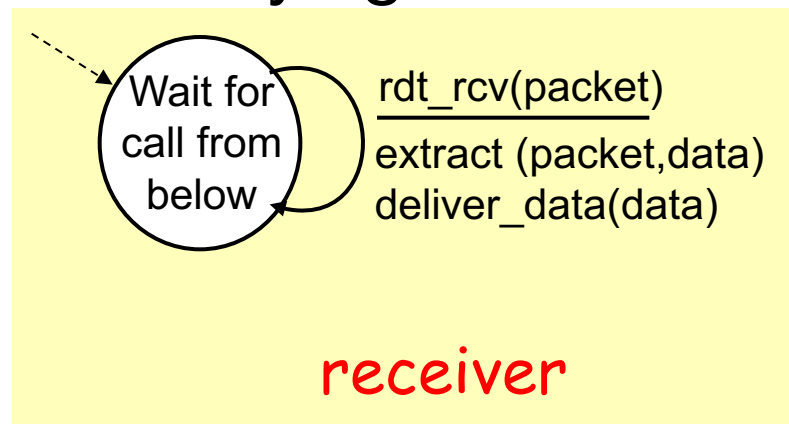
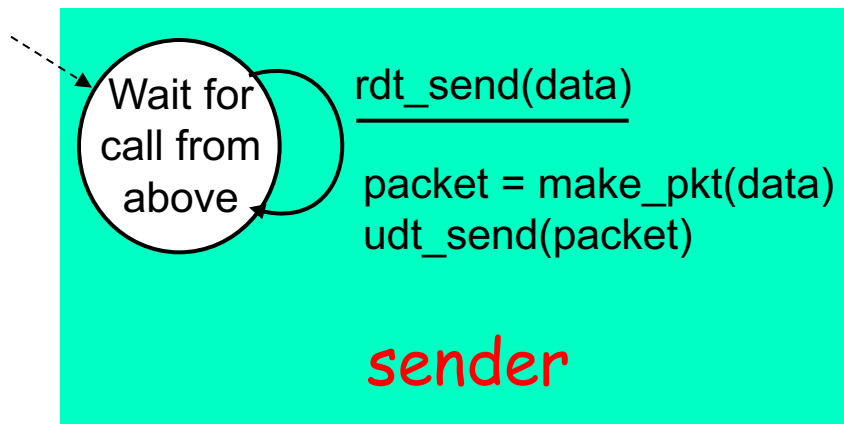


Very “simple” state diagram of TCP



Rdt1.0: reliable transfer over a reliable channel

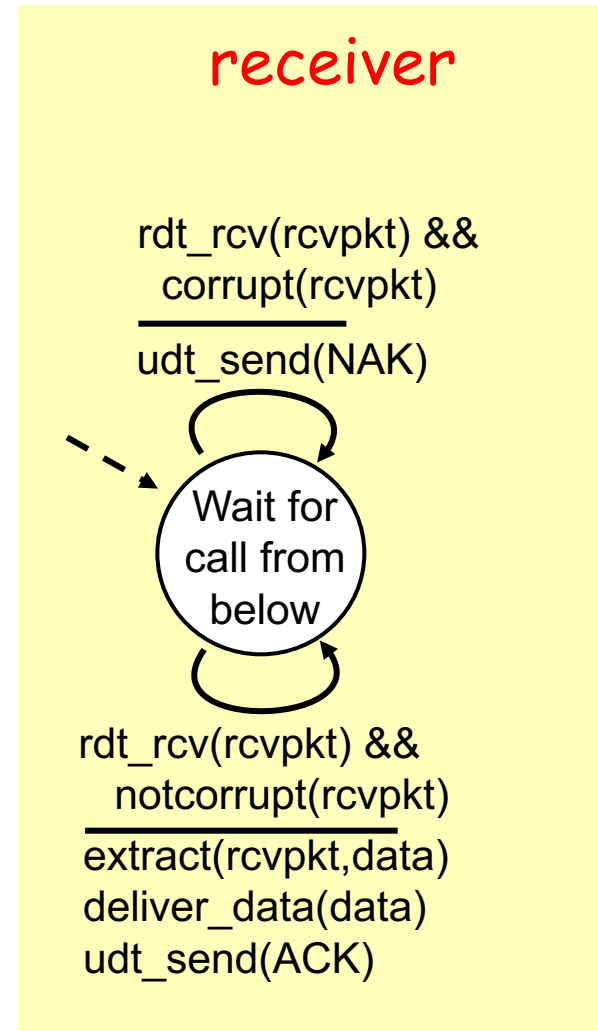
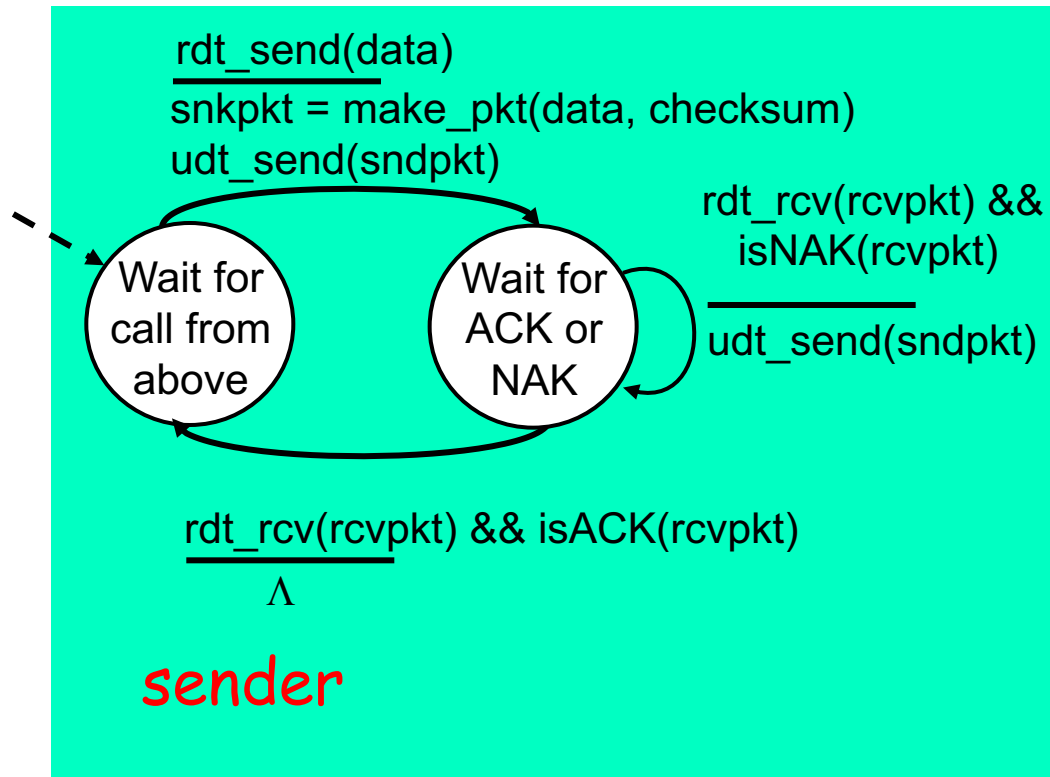
- ◆ underlying channel perfectly reliable
 - no bit errors
 - no packet losses
- ◆ separate FSMs for sender, receiver:
 - sender sends data into underlying channel
 - receiver read data from underlying channel



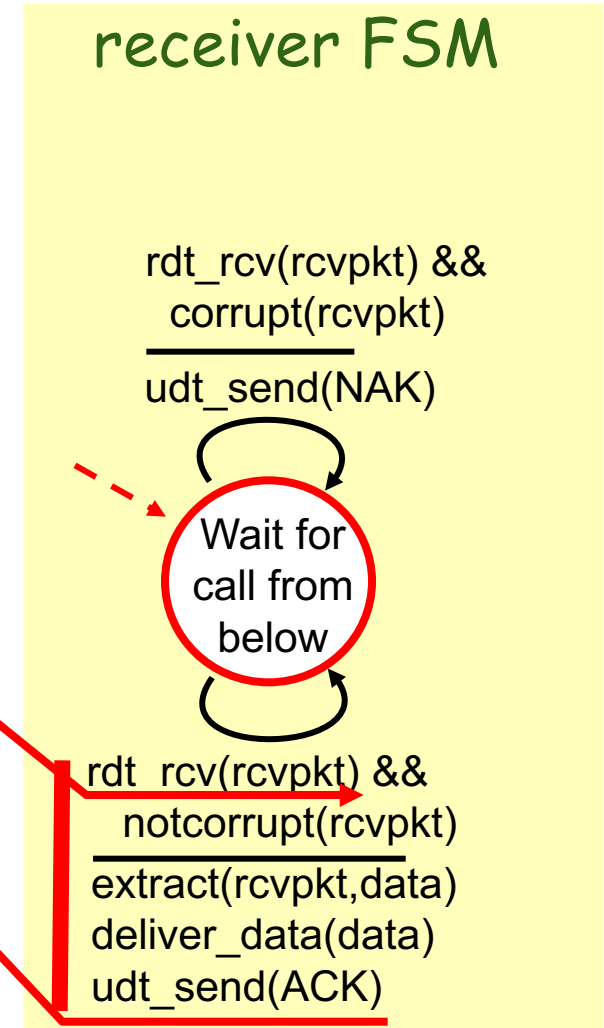
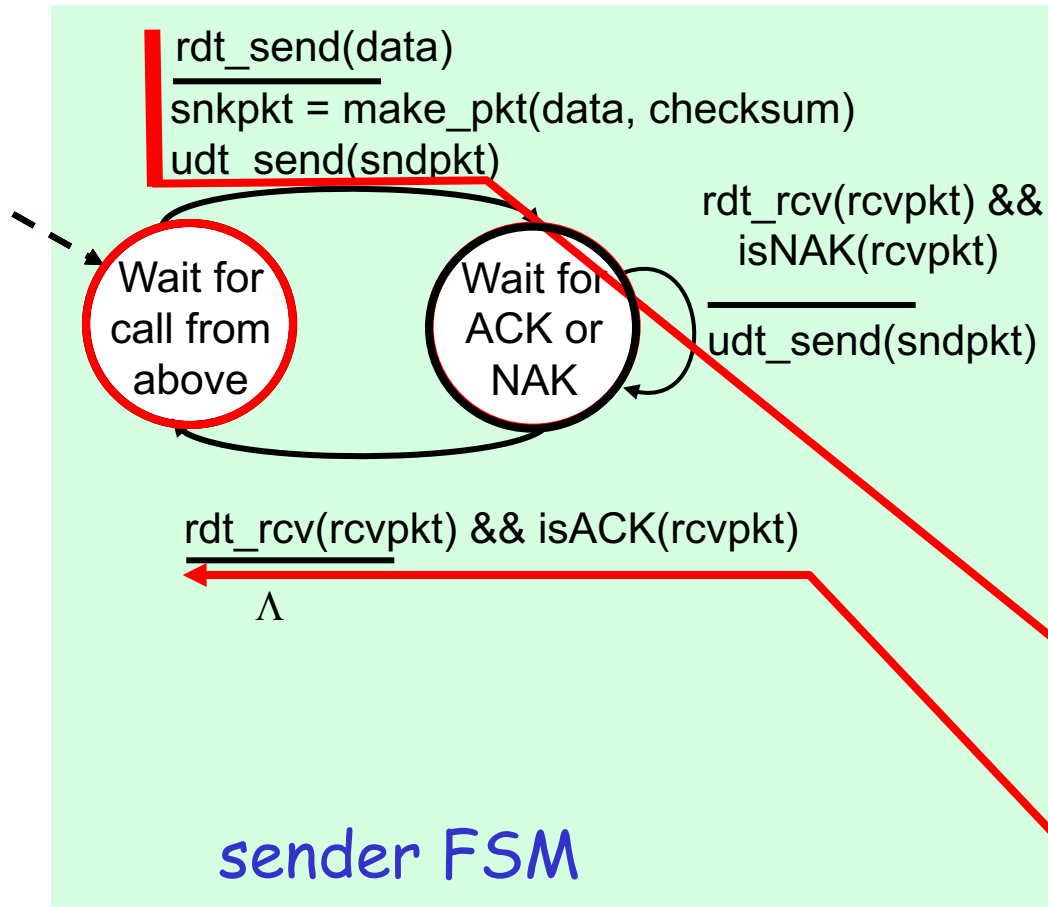
Rdt2.0: channel with bit errors

- ◆ underlying channel may flip bits in packet
 - checksum to detect bit errors
- ◆ *the question: how to recover from errors:*
 - *acknowledgements (ACKs)*: receiver explicitly tells sender that pkt received OK
 - *negative acknowledgements (NAKs)*: receiver explicitly tells sender that pkt had errors
 - sender retransmits pkt on receipt of NAK
- ◆ new mechanisms in **rdt2.0** (beyond **rdt1.0**):
 - error detection
 - receiver feedback: control msgs (ACK,NAK) rcvr → sender

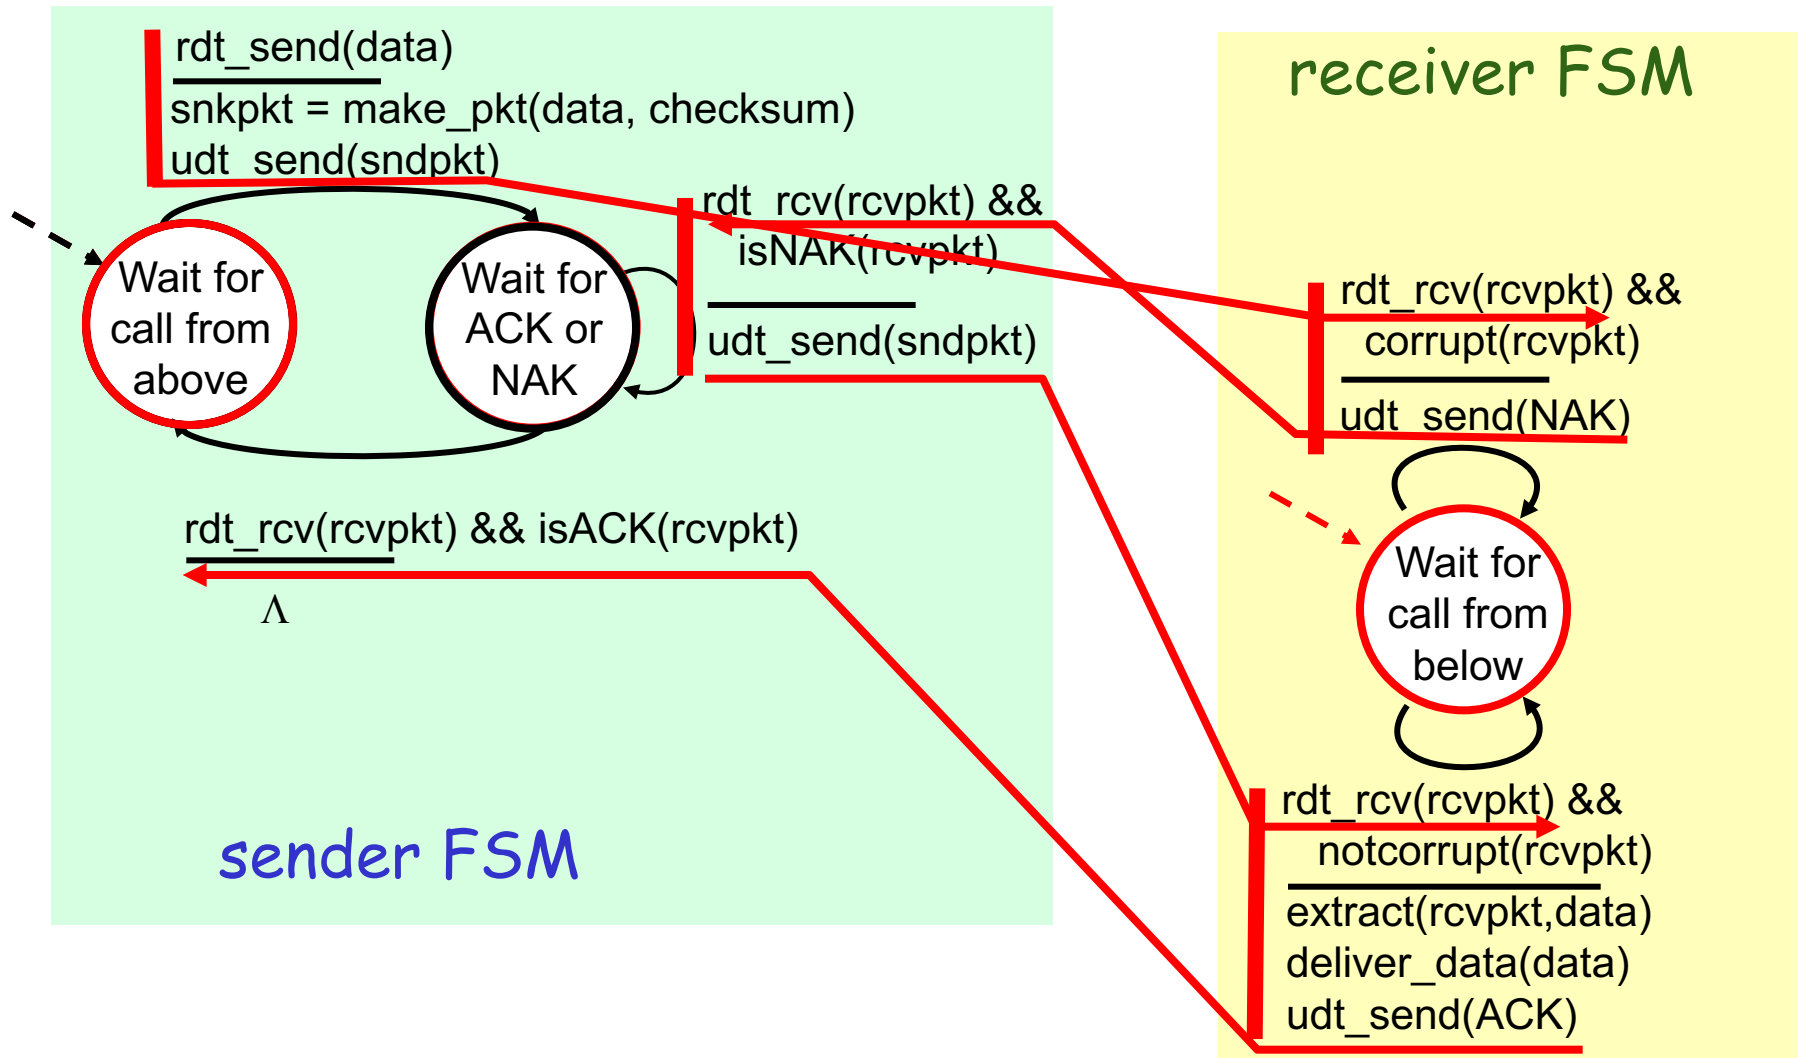
rdt2.0: FSM specification



rdt2.0: operation with no errors



rdt2.0: error scenario



rdt2.0 has a fatal flaw!

What happens if ACK/NAK corrupted?

- ◆ sender doesn't know what happened at receiver!
 - ◆ can't just retransmit: possible duplicate
- Need a way to detect duplicate

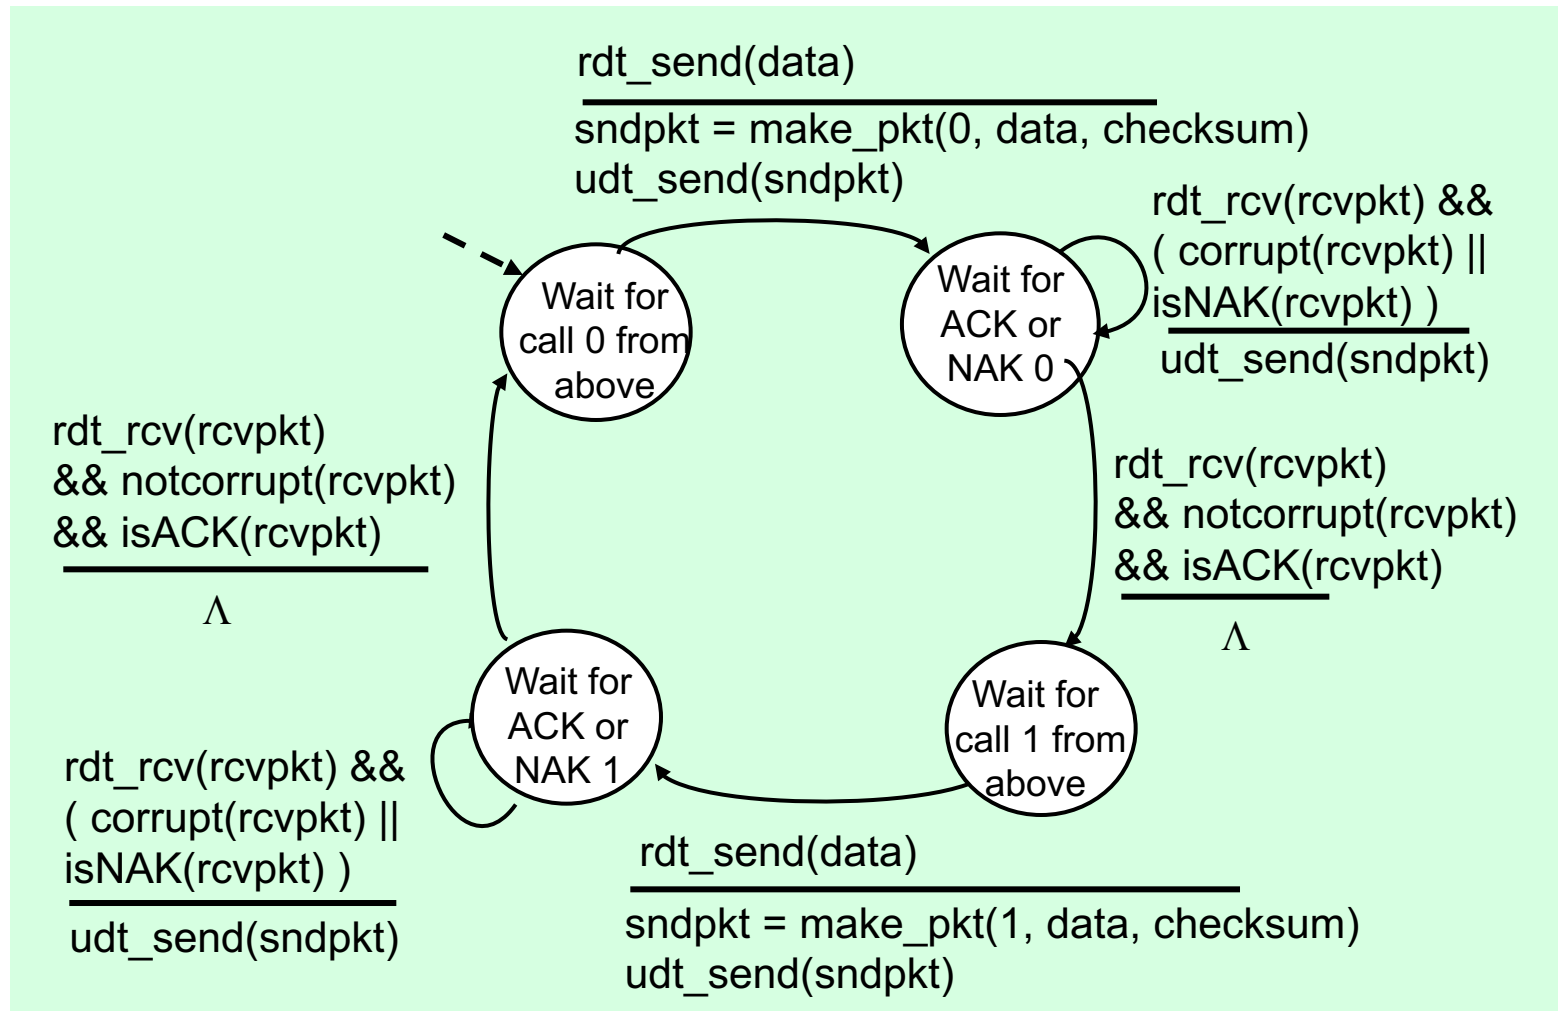
Handling duplicates:

- ◆ sender retransmits current pkt if ACK/NAK garbled
- ◆ sender adds *sequence number* to each pkt
- ◆ receiver discards duplicate pkt

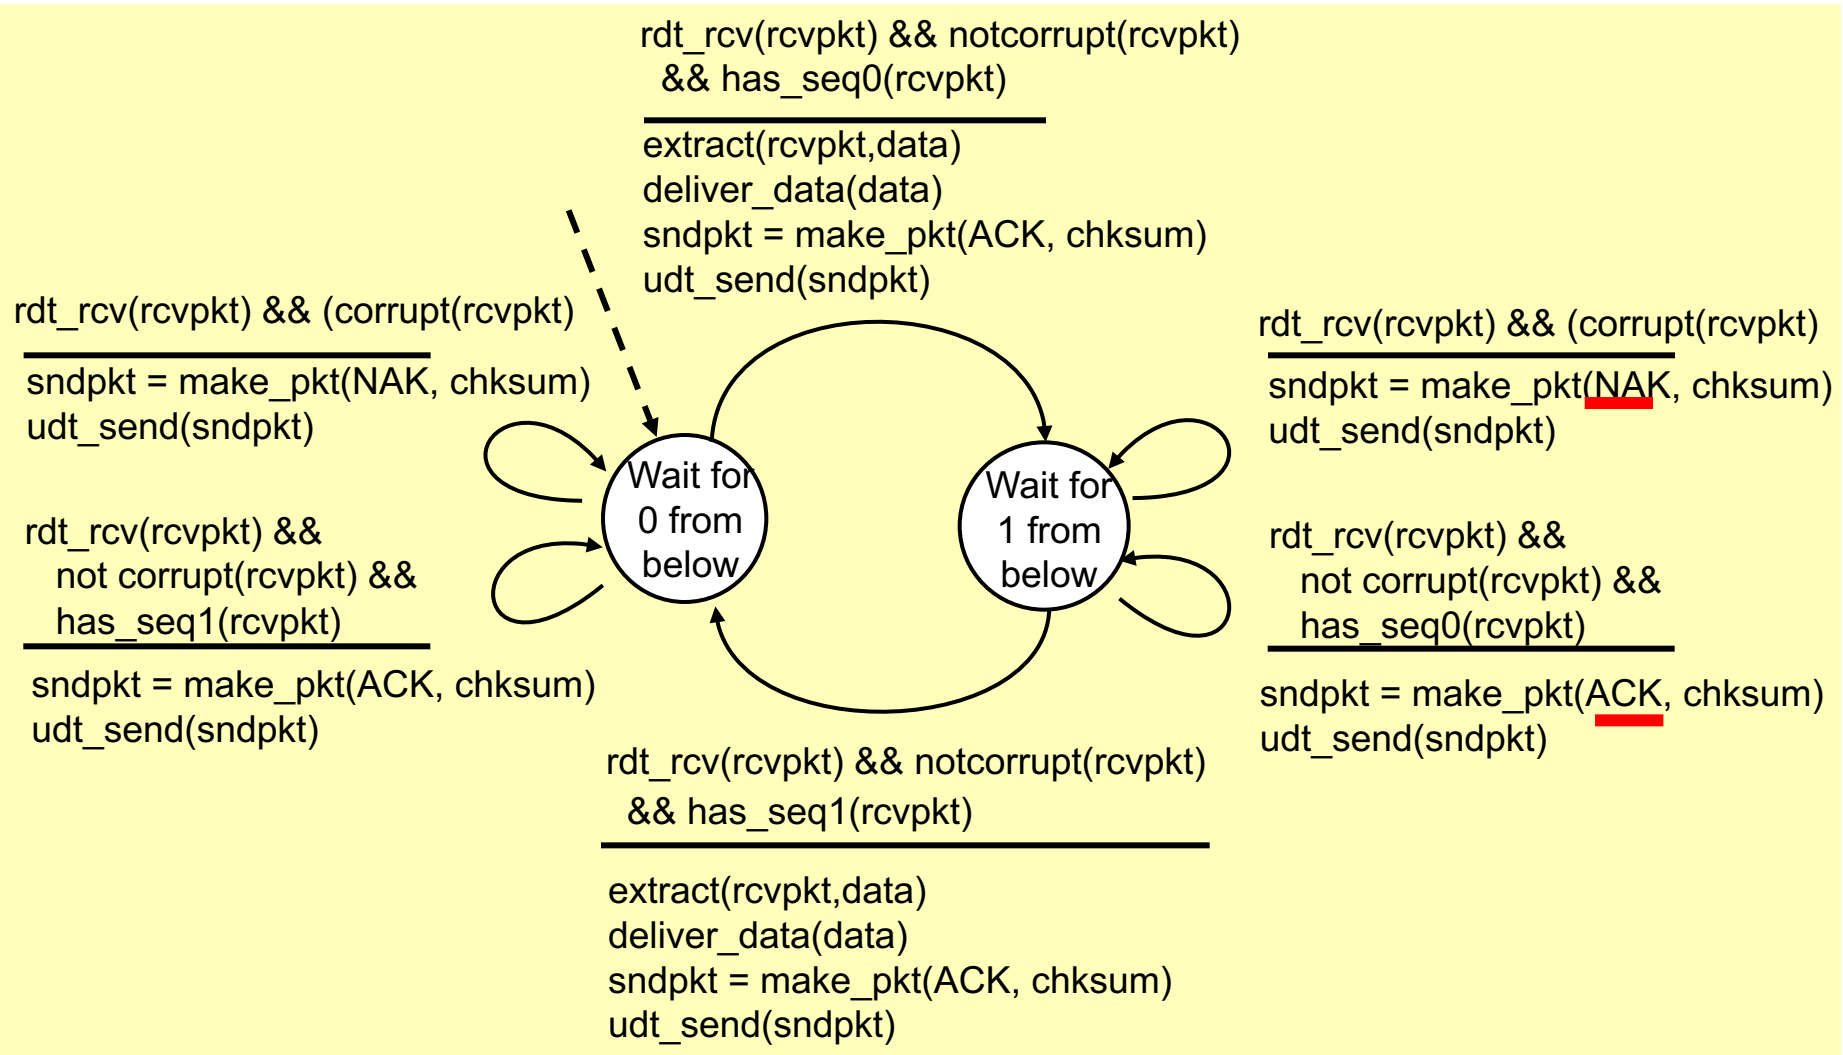
stop and wait

Sender sends one packet, then waits for receiver's response

rdt2.1: sender, handles garbled ACK/NAKs



rdt2.1: receiver, handles garbled ACK/NAKs



rdt2.1: discussion

Sender:

- ◆ seq # added to pkt
- ◆ two seq. #'s (0,1) will suffice. Why?
- ◆ must check if received ACK/NAK corrupted
- ◆ twice as many states
 - state must “remember” whether “current” pkt has 0 or 1 seq. #

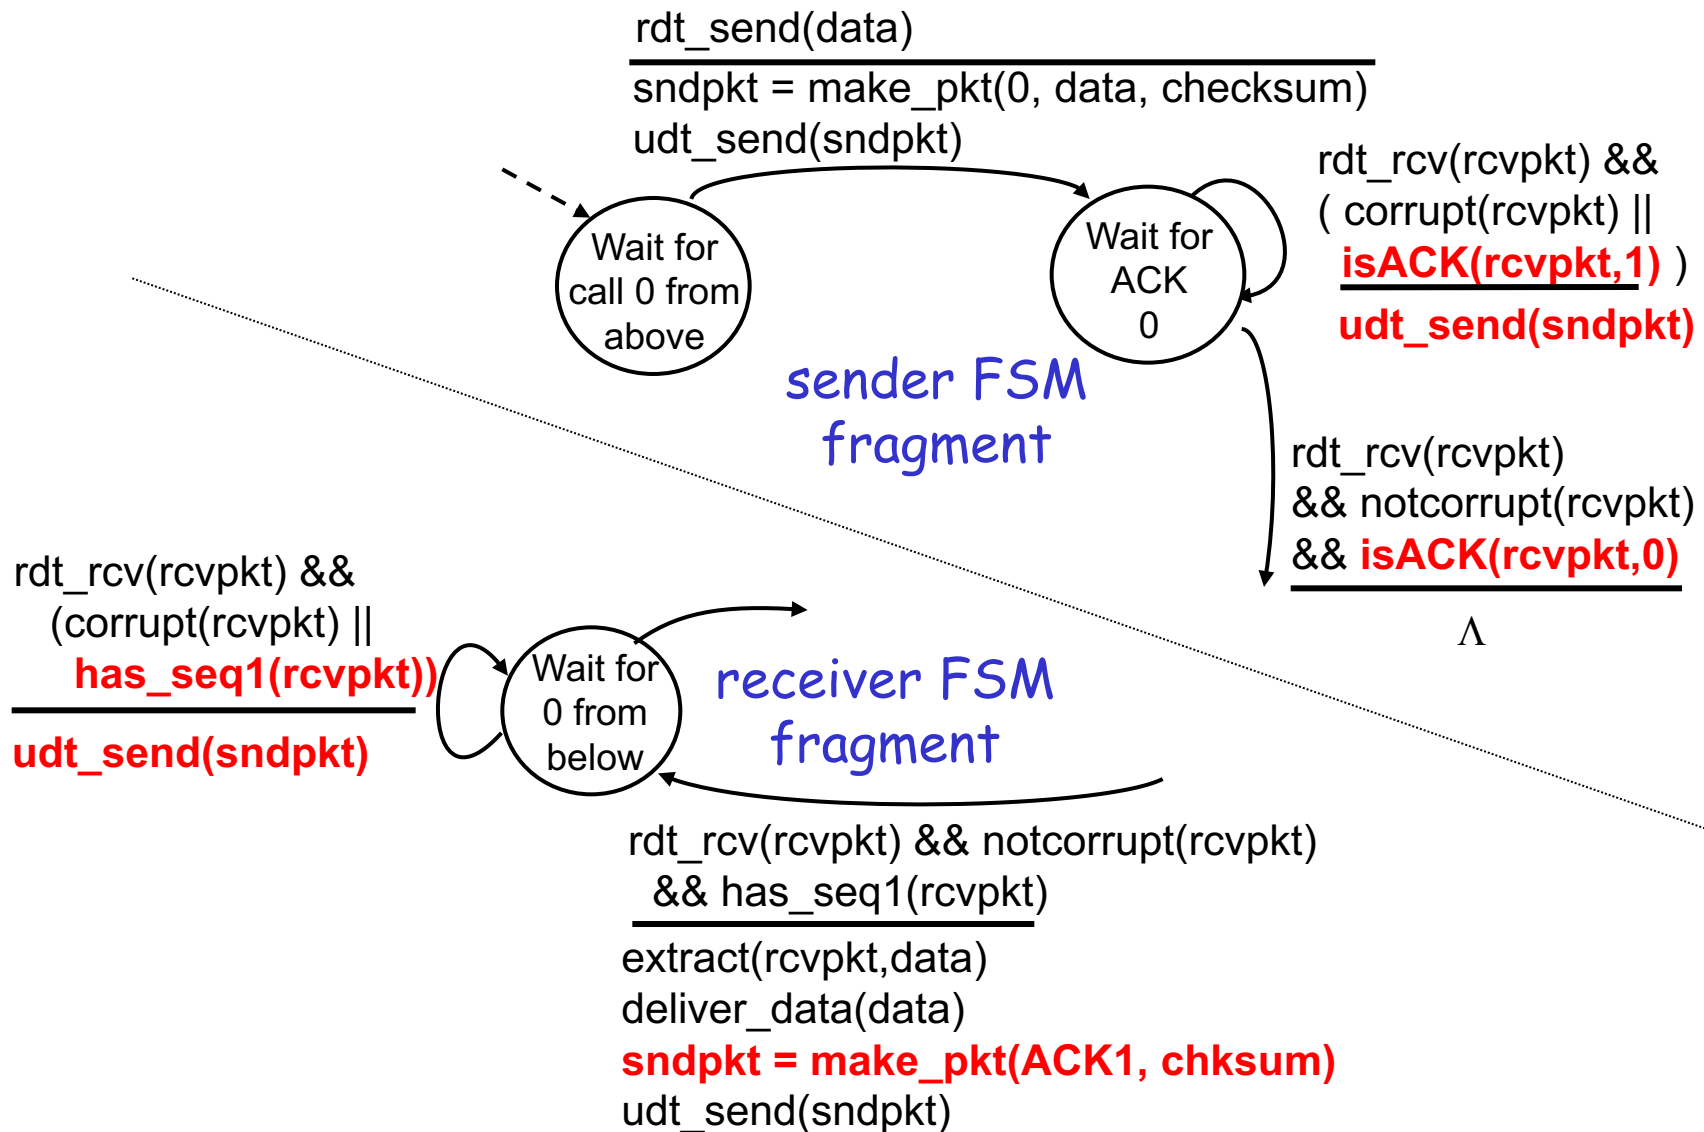
Receiver:

- ◆ must check if received packet is duplicate
 - state indicates whether 0 or 1 is expected pkt seq #
- ◆ note: receiver *cannot* know if its last ACK/NAK received OK at sender

rdt2.2: a NAK-free protocol

- ◆ same functionality as rdt2.1, using ACKs only
- ◆ instead of NAK, receiver sends ACK for last pkt received OK
 - receiver must *explicitly* include seq # of pkt being ACKed
- ◆ duplicate ACK at sender results in same action as NAK: *retransmit current pkt*

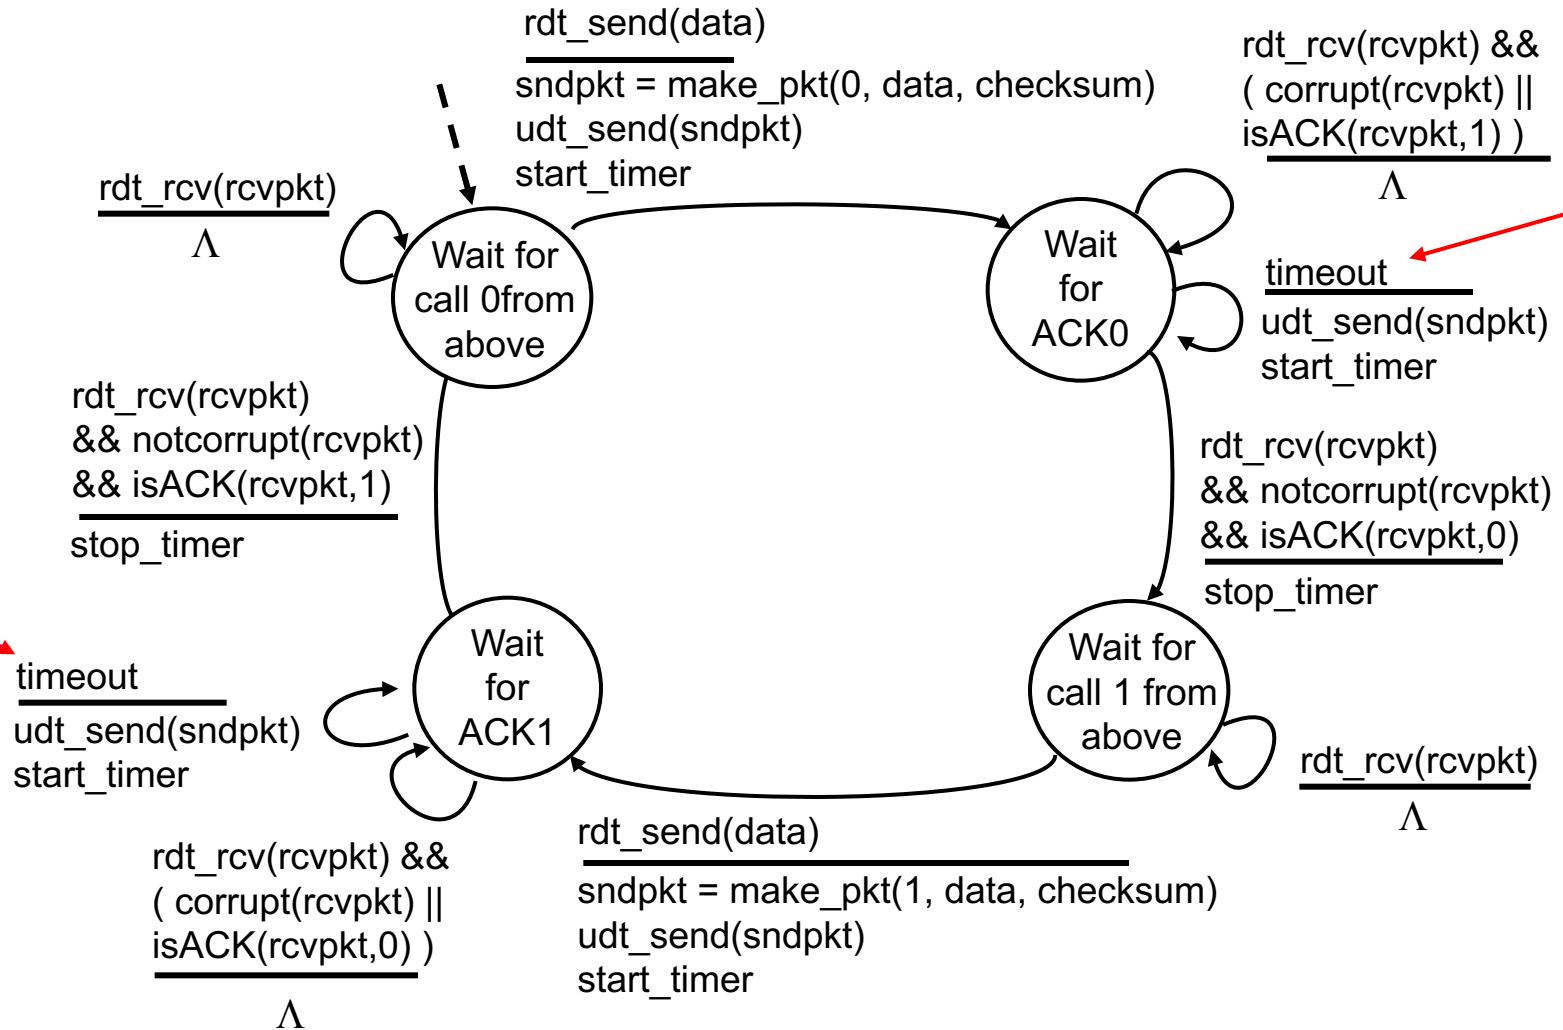
rdt2.2: sender, receiver fragments



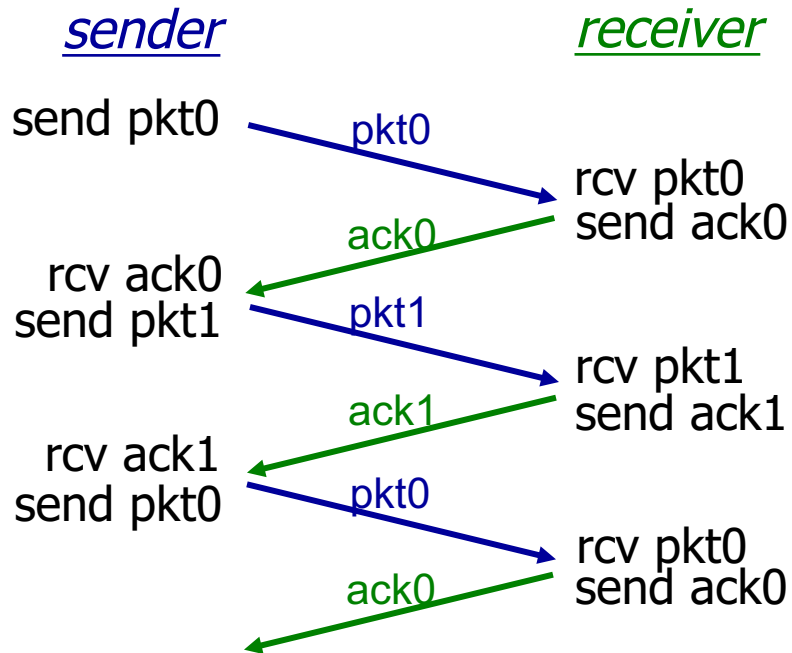
rdt3.0: channel with bit errors & packet loss

- ◆ After sending out a packet, the sender waits for **ACK** from receiver
 - Set up a **retransmission timer**
- ◆ When the timer expires: retransmits the packet
- ◆ In case the packet (or ACK) just delayed but not lost
 - Retransmitted packet will be a duplicate
 - **Sequence number** can detect this


rdt3.0 sender



rdt3.0 in action

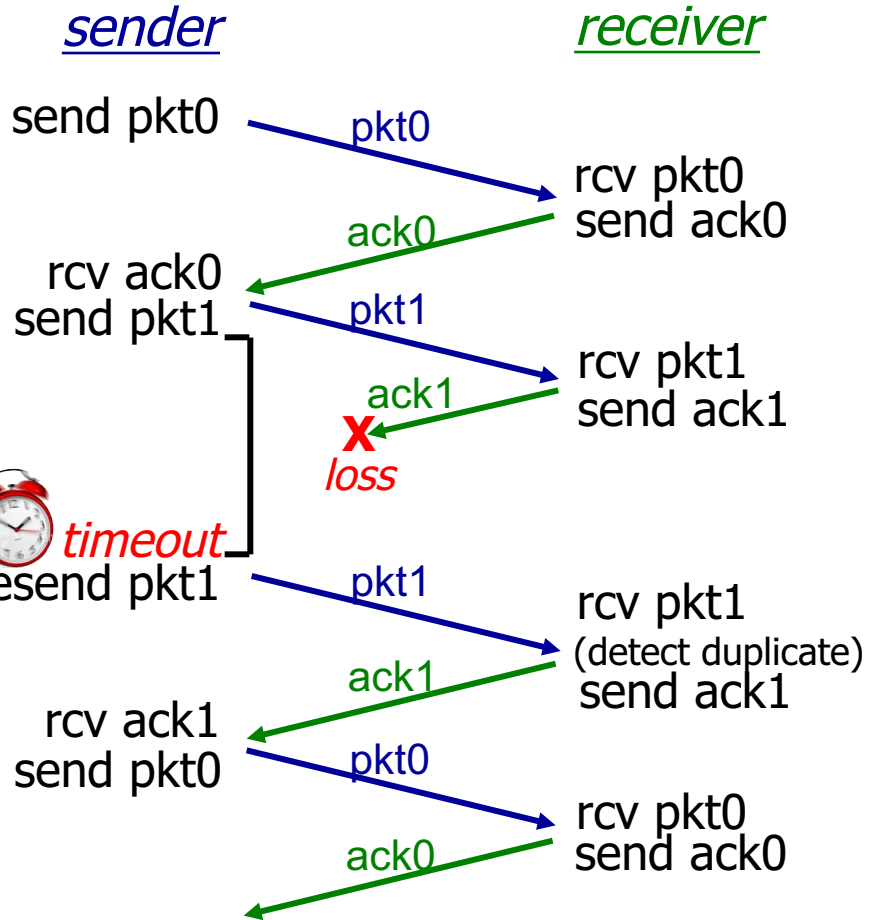


(a) no loss




(b) packet loss

rdt3.0 in action



(c) ACK loss

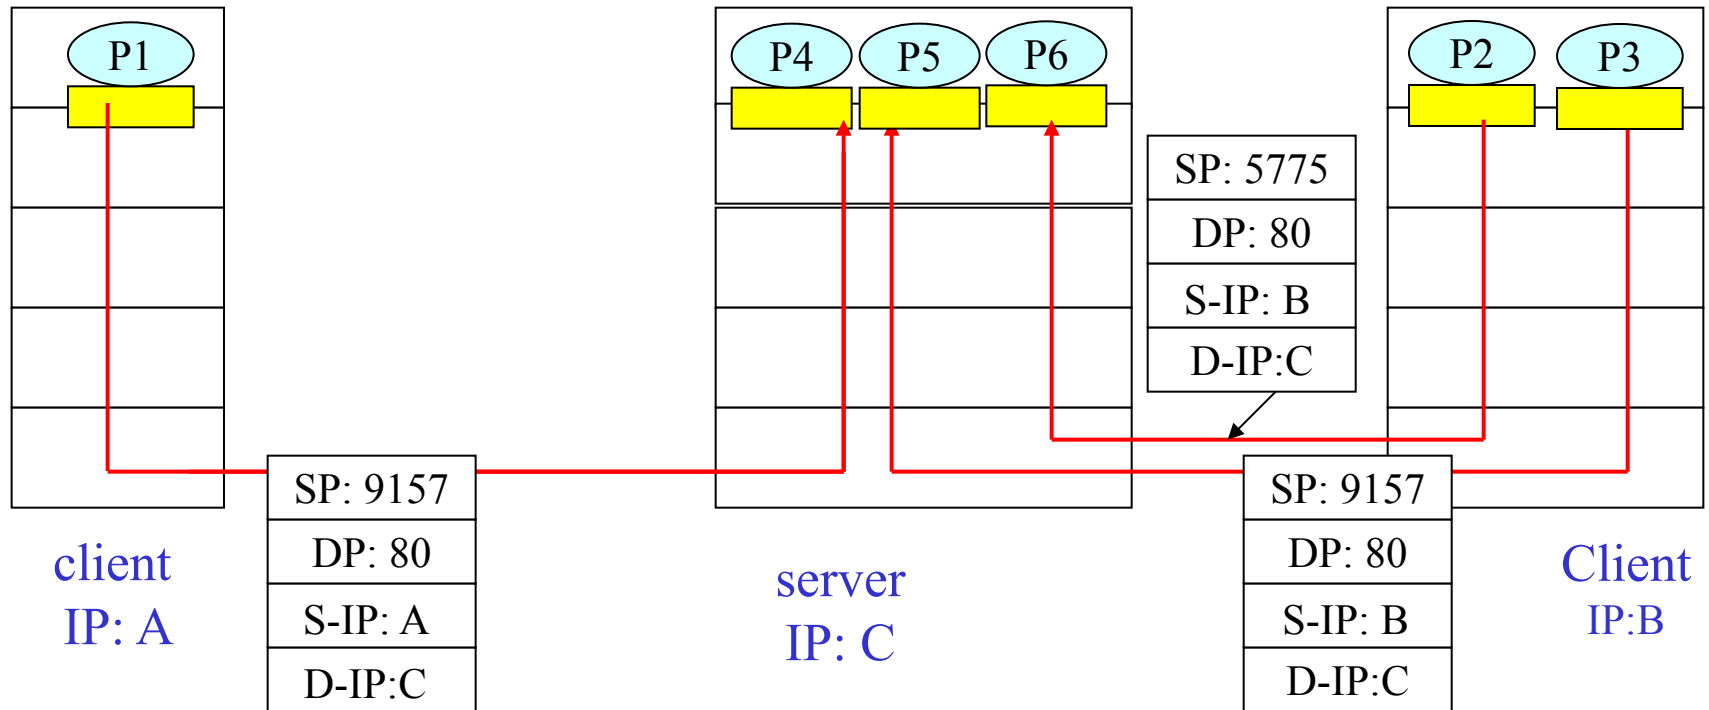


(d) premature timeout due to delayed ACK: duplicate transmissions

Connection-oriented Demultiplex

- ◆ a TCP socket identified by 4-tuple:
 - source IP address
 - source port number
 - dest IP address
 - dest port number
- ◆ receiving host uses all four values to direct segment to appropriate socket
- ◆ A server host may support many simultaneous TCP sockets:
 - each socket identified by its own 4-tuple
- ◆ e.g. a web server creates separate sockets for each connecting client

Connection-oriented demux (cont)



A server process can tell apart

- ◆ data from different hosts by IP addresses
- ◆ Data from the same host but different processes by source port numbers

Multiplexing/demultiplexing

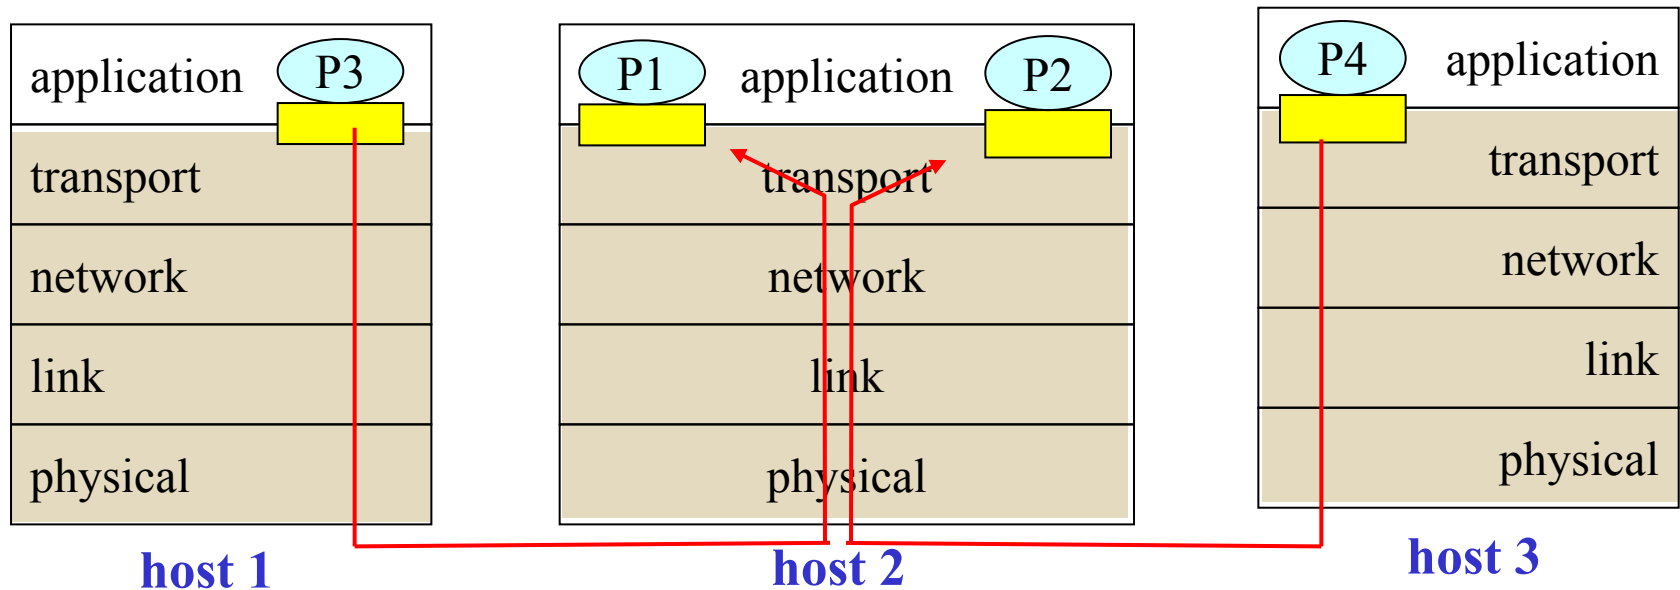
Multiplexing at sender:

gathering data from sockets,
enveloping data with header
(used for demultiplexing later)

Demultiplexing at receiver:

delivering received segments
to correct socket

 = socket  = process



Each process is identified by IP address and port#