

# Distance Vector Algorithm

Recall Link-State algorithm:

- ◆ Each node knows the complete topology graph with link costs
- ◆ Each node calculate the shortest path to all other nodes

For Distance-Vector algorithm:


- ◆ each node know *only* needs from each direct neighbor its list of distances to all destinations
- ◆ Each node computes shortest path based on the input from all its neighbors

# Distance Vector Equation

Define:  $D_x(y) :=$  cost of best path from  $x$  to  $y$

Then  $D_x(y) = \min \{c(x,v) + D_v(y)\}$

- where min is taken over *all* neighbors  $v$  of  $x$



$$\begin{aligned} D_A(F) &= \min \{c(A,B) + D_B(F), \\ &\quad c(A,D) + D_D(F), \\ &\quad c(A,C) + D_C(F)\} \\ &= \min \{2 + 5, \\ &\quad 1 + 3, \\ &\quad 5 + 3\} = 4 \end{aligned}$$

Node leading to shortest path is  
next hop → forwarding table

# Distance Vector: what a node does

- ◆ Node  $x$  knows link cost to neighbor  $v$ :  $c(x,v)$
- ◆ Node  $x$  maintains  $\mathbf{D}_x = [D_x(y): y \in N]$ 
  - $D_x(y)$  = estimate of least cost from  $x$  to  $y$
- ◆ Node  $x$  sends the distance vector,  $\mathbf{D}_x$ , to all its neighbors
- ◆ Node  $x$  receives  $\mathbf{D}_v$  from each neighbor  $v$ , then calculate  $D'_x(y) = \min \{c(x,v) + D_v(y)\}$ 
  - If  $D'_x(y) < D_x(y)$ :
    - $D_x(y) = D'_x(y)$
    - next hop to  $y = v$
    - Send out the updated  $\mathbf{D}_x$

# Distance Vector Protocol

## Iterative, asynchronous:

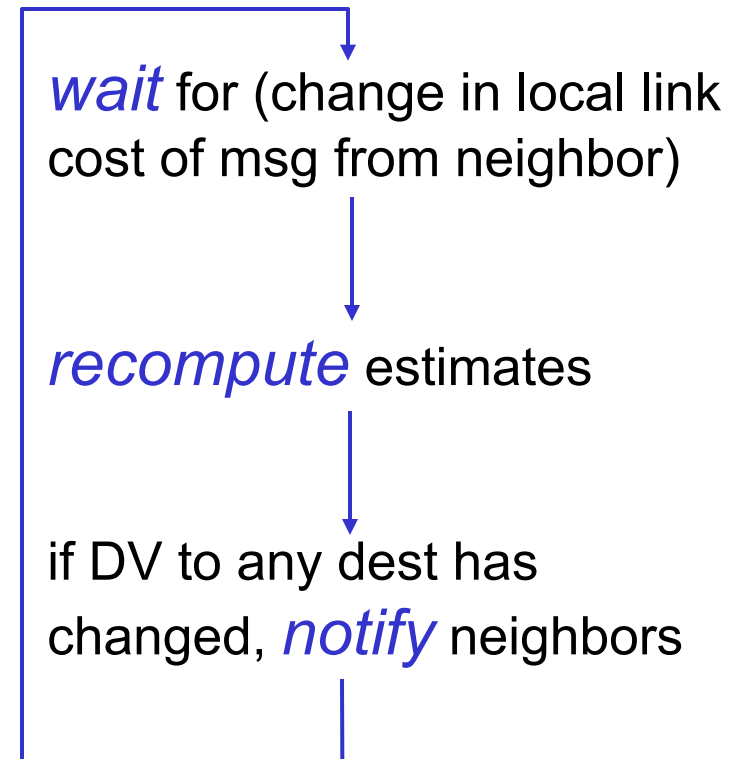
- ◆ each local iteration caused by:
  - local link cost change
  - DV update message from neighbor
- ◆ continues until no nodes exchange info.

## Distributed:


- ◆ each node notifies neighbors *only* when its DV changes
  - neighbors then notify their neighbors if necessary

**asynchronous:** nodes need *not* exchange info/iterate in lock step

## Each node:



# Distance Table: example



cost to destination via

$D^E()$	A	B	D
A	1	14	5

destination

$$D^E(A, B) = c(E, B) + \min_w \{D^B(A, w)\} \\ = 8 + 6 = 14$$

$$D^E(A, D) = c(E, D) + \min_w \{D^D(A, w)\} \\ = 2 + 3 = 5$$

$$D^E(C, D) = c(E, D) + \min_w \{D^D(C, w)\} \\ = 2 + 2 = 4$$

**Row:** for each possible destination  
**Column:** for each directly-attached neighbor node

# Routing table produces forwarding table

cost to destination via


$D^E()$	A	B	D
A	1	14	5
B	7	8	5
C	6	9	4
D	4	11	2

destination

Next hop


A	A
B	D
C	D
D	D

destination

E's routing table  E's forwarding table

# Count-To-Infinity Problem


- ◆ Assume we use hop count as metric
  - A uses B to reach D with cost 3
  - B uses C to reach D with cost 2
  - C reaches D with cost 1




- ◆ Suppose link between C and D breaks
  - C switches to B, increase its cost to  $B's + 1 = 3$

# Count-To-Infinity Problem (cont.)

- ◆ B's path cost is now 4
  - A has not realized what has happened yet



- ◆ Then, A's and C's cost are now 5




- ◆ B's path cost is changed to 6
  - Cycle repeats while “counting to infinity”




# Routing Loops

- ◆ In this cases, the packets with destination of D at router A:
  - Go to router B
  - Then go to router C
  - Then go back to router B



# Split Horizon


- ◆ In Split Horizon: B does not tell C that B can reach D
  - So C does not know that B can reach D



- ◆ Once C-D link breaks: C would not switch to go through B to reach D

# Split Horizon --- Might Not Work

- ◆ Split Horizon doesn't eliminate loops in all cases
- ◆ Suppose the link between C and D breaks




1. A and B do not tell C their distance to D
2. After C–D failure, A learns that B can reach D, so sends new route to C
3. C sends route learned from A to B
4. B sends route learned from C to A
5. A sends route learned from B to C

Routing loop still exists

# Split Horizon with poison reverse

- ◆ If B goes through C to reach D :
  - B tells C that its (B's) distance to D is *infinite* (so C never attempts to reach D via B)



1. A and B tell C that their distance to D is infinite
2. When link C-D fails, C realized that it lost reachability to D
3. C sends to A and B:  $D_D = \text{infinite}$

# Comparison of LS and DV algorithms

- ◆ Performance measure: Message overhead, time to convergence
- ◆ distance vector:
  - distribute to neighbors the distances to all destinations
    - Each update msg can be large in size, but travels over one link
  - each node only knows *distances* to other destinations
- ◆ link state
  - Broadcast to entire net one's distance to all neighbors
    - Each update msg is small in size, but travels over all links in the network
  - each node knows entire topology


# what happens if a router malfunctions?

## ◆ Link-state

- A node can advertise incorrect *link* cost
- each node computes its *own* table

## ◆ Distance vector

- A node can advertise incorrect *path* cost
- one node's distance-list is used by its neighbors for their own routing selection



**Node-D: “I have 0 cost to all other nodes”**

Link-State:

- updates from A & B: not connected to D
- Updates from C & E: cost not 0

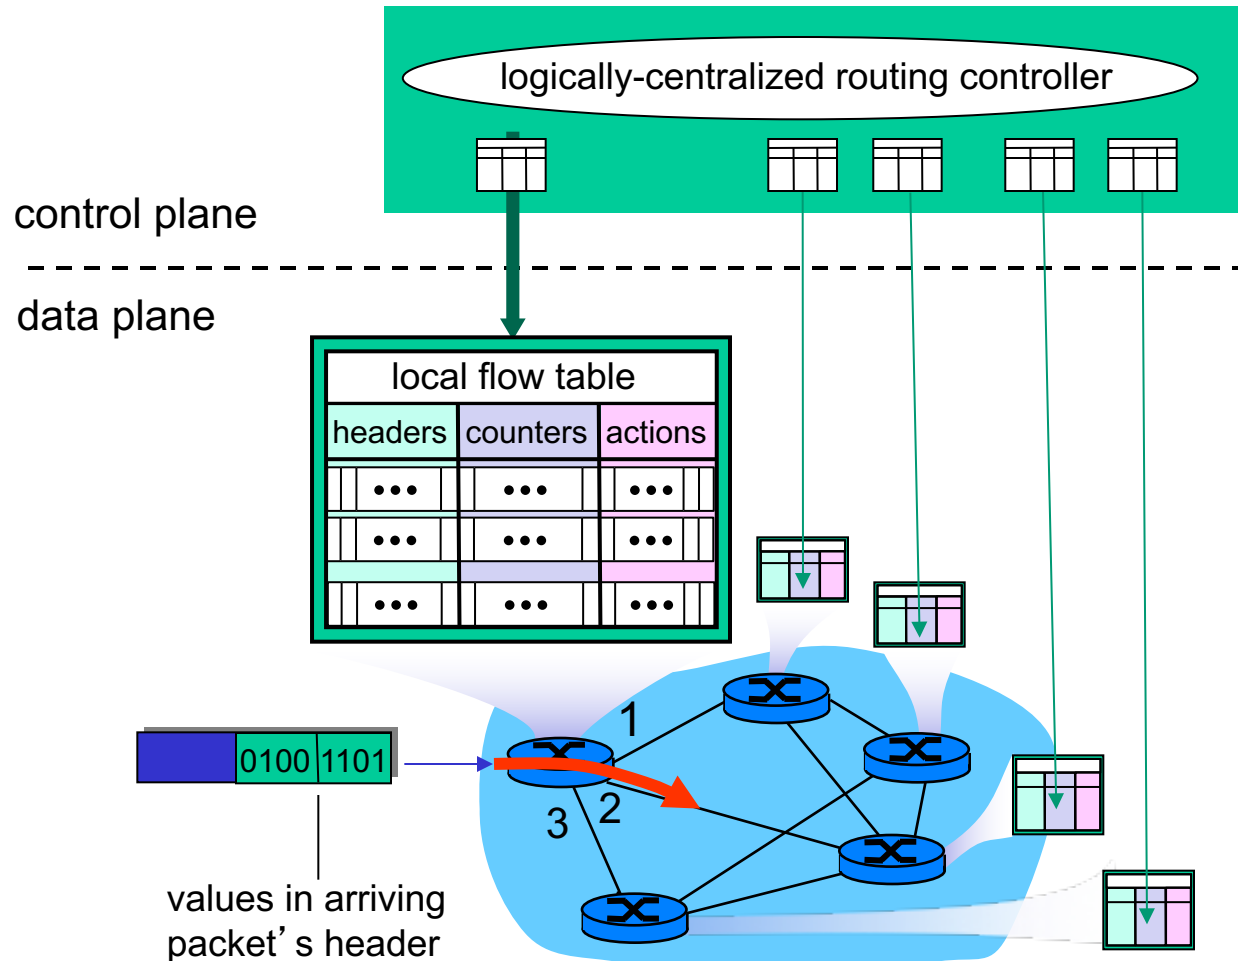
Distance-Vector:

- other nodes do not have info to verify

# **Software Defined Networking (SDN)**

# Generalized Forwarding and SDN


Each router contains a *flow table* that is computed and distributed by a *logically centralized routing controller*





# OpenFlow data plane abstraction


- ◆ *flow*: defined by header fields (may across multiple protocol layers)
- ◆ generalized forwarding: simple packet-handling rules
  - *Pattern*: match values in packet header fields
  - *Actions: for matched packet*: drop, forward, modify, matched packet or send matched packet to controller
  - *Priority*: disambiguate overlapping patterns
  - *Counters*: #bytes and #packets



*Flow table in a router (computed and distributed by controller) define router's match+action rules*

# OpenFlow data plane abstraction

- ◆ *flow*: defined by header fields
- ◆ generalized forwarding: simple packet-handling rules
  - *Pattern*: match values in packet header fields
  - *Actions: for matched packet*: drop, forward, modify, matched packet or send matched packet to controller
  - *Priority*: disambiguate overlapping patterns
  - *Counters*: #bytes and #packets




1. src=1.2.\*.\*, dest=3.4.5.\* → drop \* : wildcard
2. src = \*.\*.\*.\*, dest=3.4.\*.\* → forward(2)
3. src=10.1.2.3, dest=\*.\*.\*.\* → send to controller

# OpenFlow example

*Example:* datagrams from hosts **h5** and **h6** should be sent to h3 or h4, via s1 and from there to s2

match	action
IP Src = 10.3.*.* IP Dst = 10.2.*.*	forward(3)



match	action
ingress port = 1 IP Src = 10.3.*.* IP Dst = 10.2.*.*	forward(4)

match	action
ingress port = 2 IP Dst = 10.2.0.3	forward(3)
ingress port = 2 IP Dst = 10.2.0.4	forward(4)

# Routing: What we have learned so far

- ◆ Link-state routing (Dijkstra) algorithm: **each node** computes the shortest paths to all the other nodes based on the complete *topology map*
- ◆ Distance Vector (Bellman-Ford) routing algorithm: **each node** computes the shortest paths to all the other nodes based on its *neighbors distance to all destinations*
- ◆ Today: routing protocols to implement them
  - Distance vector:  $D_x(y) = \min \{c(x,v) + D_v(y)\}$   
the protocol: a node must send *its distance to all destinations* to all its neighbors
  - Link-state  $\Rightarrow$  the protocol must deliver the *topology map* to all the nodes in the network

# What else a routing protocol must do


- ◆ Recover from packet losses in routing data delivery
- ◆ Monitor link and neighbor nodes status
  - Once a failure is detected, needs to inform the rest of the network quickly (directly or indirectly)
- ◆ Flush obsolete information out of the system

# OSPF: OPEN SHORT PATH FIRST

<https://tools.ietf.org/html/rfc2328>

# OSPF: Building a complete network graph using Link State

- ◆ Every node broadcasts a piece of the topology graph
- ◆ assemble all the pieces together, you get the complete graph



Then each node carries out its own routing calculation *independently*

# OSPF (Open Shortest Path First)

- ◆ Given: each node knows its directly connected neighbors & the link distance to each neighbor
- ◆ Each node periodically broadcasts its link-state to the *entire* network
  - **Delivered by raw IP packet (protocol ID = 89)**

No.		Time	Source	Destination
2	⊞	1.663948	10.0.0.6	224.0.0.5
3	⊞	3.584090	10.0.0.10	224.0.0.5
4	⊞	4.894103	10.0.0.1	224.0.0.5
5	⊞	4.894132	10.0.0.5	224.0.0.5

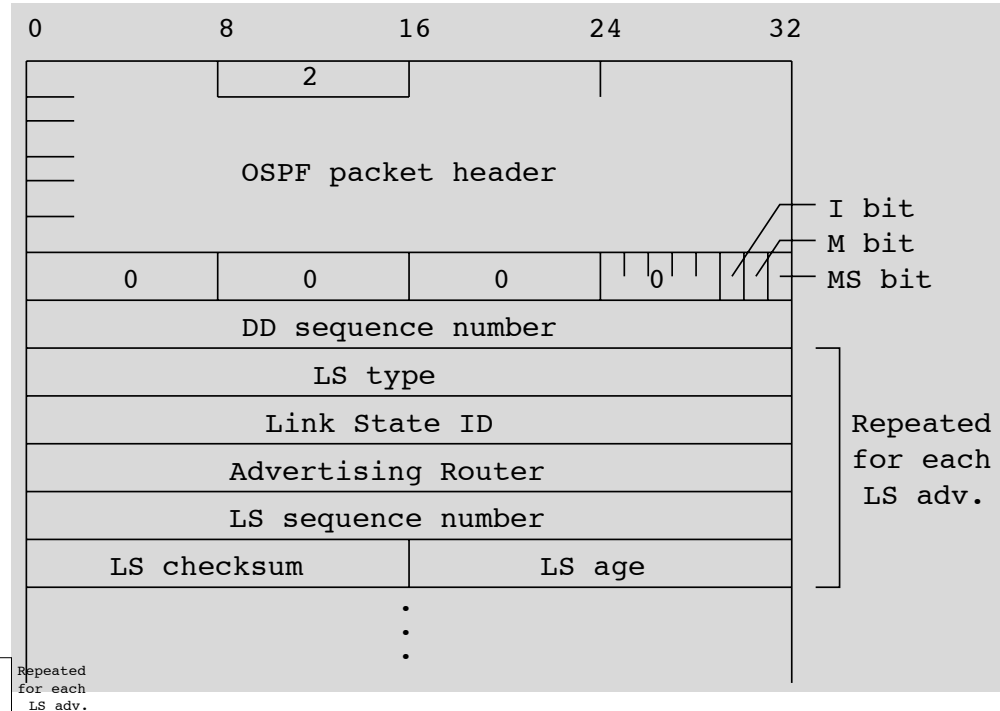


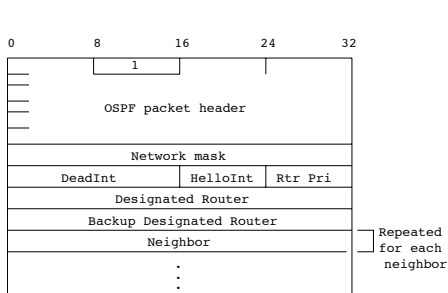
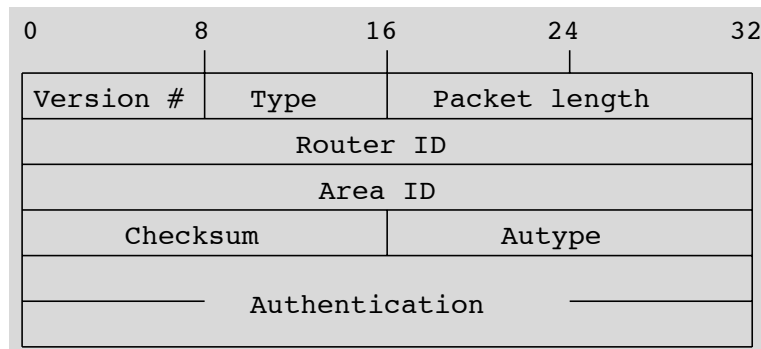
  

▷	Frame 3: 80 bytes on wire (640 bits), 80 bytes captured (640 bits)
▷	Frame Relay
▷	Internet Protocol Version 4, Src: 10.0.0.10 (10.0.0.10), Dst: 224.0.0.5 (224.0.0.5)
▽	Open Shortest Path First
▷	OSPF Header
▷	OSPF Hello Packet
▷	OSPF LLS Data Block



# OSPF (Open Shortest Path First)

- ◆ Link-State Packet (LSP): one entry per neighbor router
  - ID of the node that created the LSP
  - a list of direct neighbors, with link cost to each
  - sequence number (SEQ) for this LSP message
  - time-to-live (TTL) for information carried in this LSP



# How OSPF Works

- ◆ When neighboring routers discover each other for the first time: Exchange their link-state databases
- ◆ Link failure detection
  - Neighbor nodes send HELLO msg to each other periodically
  - Not receiving HELLO message for long enough time → **failure** → Trigger new Link State Update to neighbors
- ◆ In the absence of failure: send out update every **30 minutes**


# Link-State Routing Protocol

The routing daemon running at each node:

- ◆ Generates its own LSP periodically with increasing sequence #
- ◆ Stores most recent LSP from all other nodes
  - decrement TTL of stored LSP; discard a LSP when its TTL=0
- ◆ Process received updates to build & maintain topology graph
  - Route computation using Link-State algorithm
- ◆ Forward most recent LSPs

# Reliable Flooding of LSP

- ◆ forward each received new LSP to all neighbor nodes but the one that sent it
  - each LSP is reliably delivered over each link
  - use the sender-ID and SEQ in a LSP to detect duplicates
- ◆ LSPs sent both periodically and event-driven



Q: How many LSP msgs traverse each link in the absence of failures in an hour?


# ROUTING ON THE INTERNET

# Routing in the Internet

- ◆ So far: all routers faithfully execute the same routing protocol
  - Goal: Find best path (sequence of routers) through network from source to destination
    - Based on delay, loss, bandwidth, or other measures
- ◆ The Global Internet: interconnection of a large number of **Autonomous Systems (AS)**
  - Stub AS: end user networks (corporations, campuses)
    - Multihomed AS: stub ASes that are connected to multiple service providers
  - Transit AS: Internet service provider

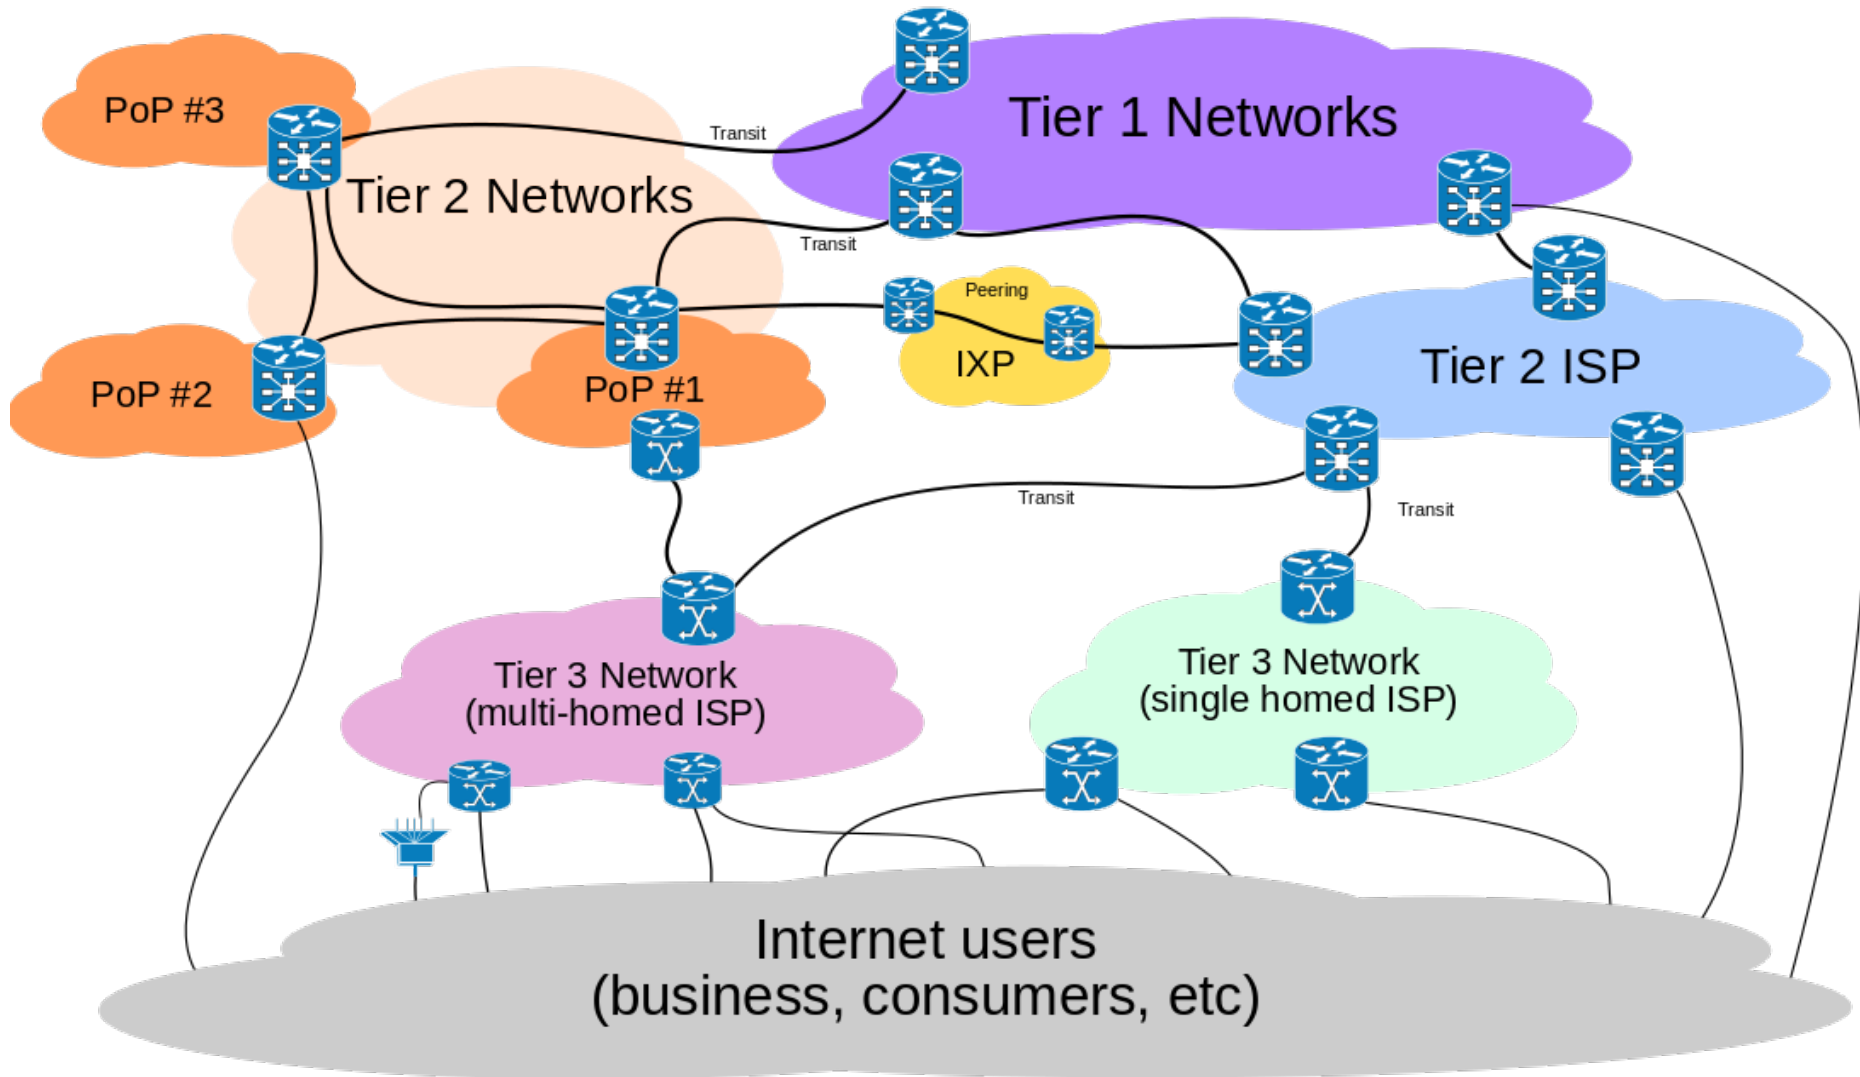
# Internet routing structure: 2-level hierarchy

- ◆ Intra-AS (within a campus, or within an ISP)
  - Intra-Domain Routing: RIP, OSPF (and a few others)
- ◆ Inter-AS (between ISPs, or between stub and transit ASes)
  - Inter-Domain Routing: BGP (Border Gateway Protocol)



# Internet Structure

FYI



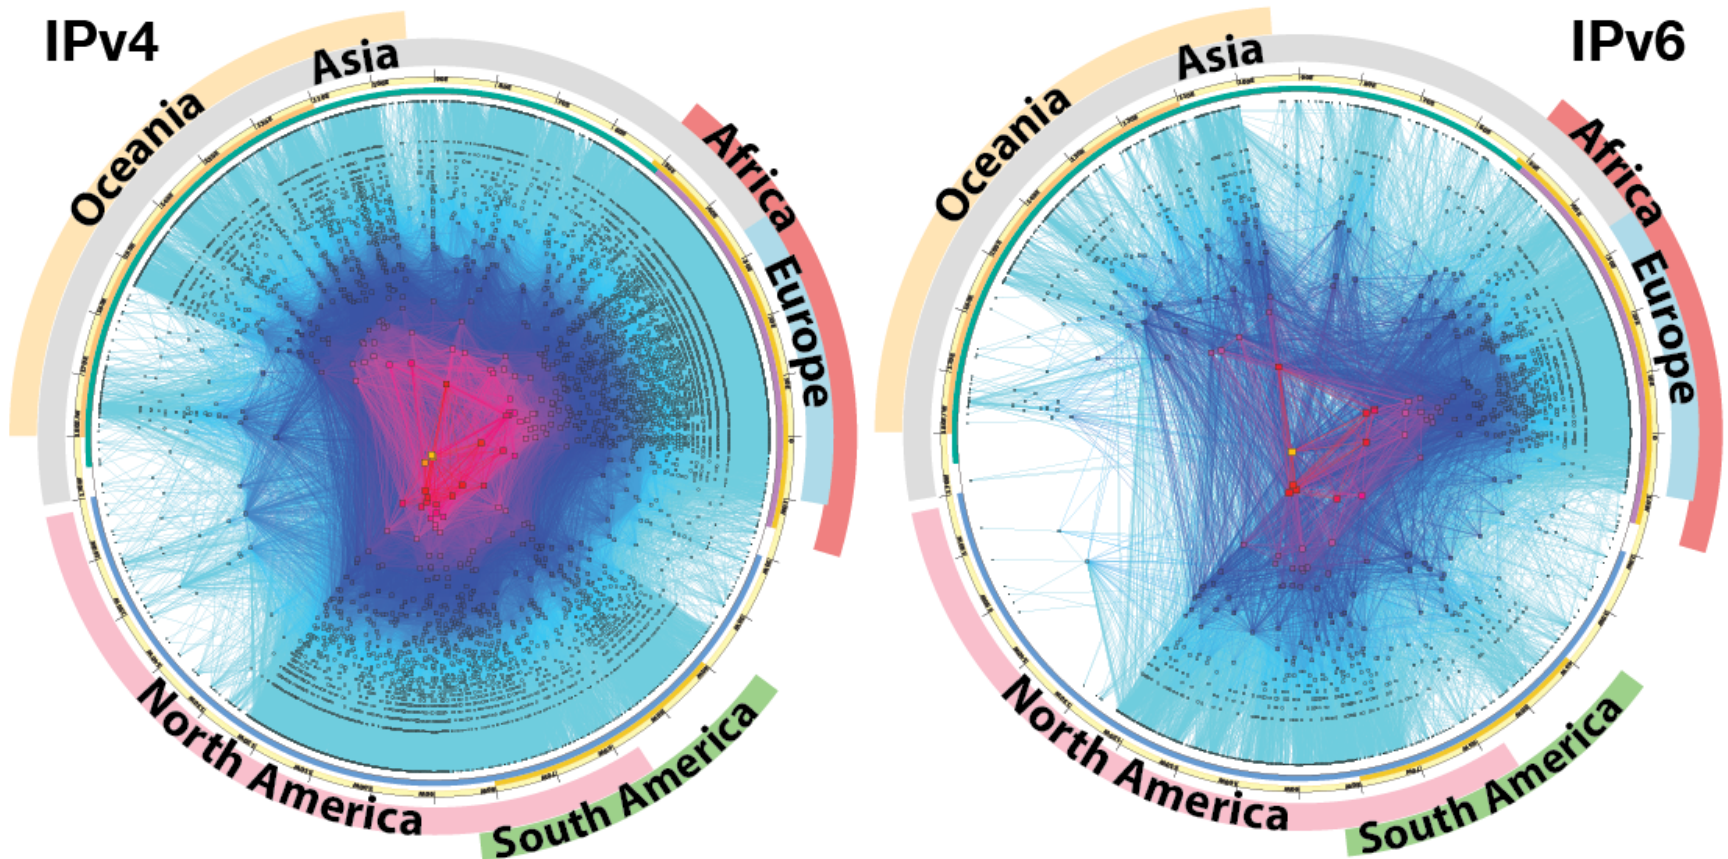


# Internet Interconnections as a Graph

FYI

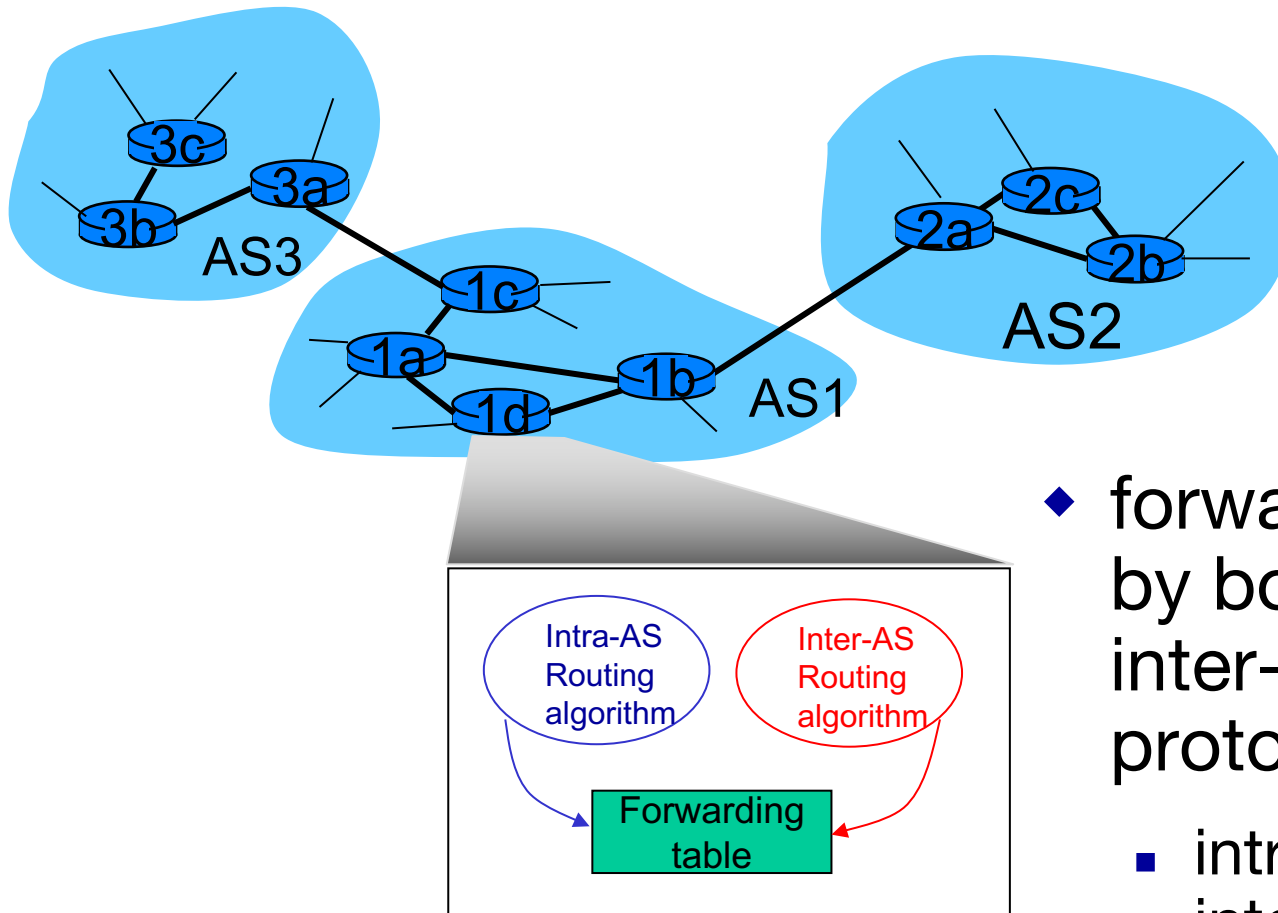
## CAIDA's IPv4 vs IPv6 AS Core AS-level Internet Graph

Archipelago July 2015



Copyright © 2015 UC Regents. All rights reserved.

# Interconnected ASes



- ◆ forwarding table filled by both intra- and inter-AS routing protocols
  - intra-AS sets entries for internal dests
  - inter-AS & intra-AS sets entries for external dests

# BGP: Border Gateway Protocol

BGP provides each AS a means to:

- ◆ Obtain each subnet reachability (= an IP address prefix) information from neighboring ASes.
- ◆ Propagate the reachability information to all routers internal to the AS.
- ◆ Determine “*good*” routes to each destination prefix based on reachability information and policy.
- ◆ advertise its own prefixes to the rest of the Internet

CalREN  
AS2513




UCLA  
AS52

# BGP basics: distributing path information


important

- ◆ 2 neighbor BGP routers establish a BGP session over a (semi-permanent) **TCP connection** to exchange routing info
  - advertising **paths** to destination network prefixes (“path vector” protocol)
- ◆ when AS3 advertises a prefix to AS1:
  - AS3 promises it will forward packets towards that prefix




# BGP basics: eBGP and iBGP

- ♦ **eBGP**: BGP session between two different ASs
  - e.g. the BGP session between 3a and 1c
- ♦ **iBGP**: BGP session between routers in the same AS
  - Router **1c** uses iBGP to distribute new prefix info (e.g. 12.3.6.0/24) to all routers in AS1
  - when router learns of new prefix, it creates entry for prefix in its forwarding table.
  - Router **1b** may re-advertise this reachability info to AS2 over 1b-to-2a eBGP session



# Path attributes and BGP routes


- ◆ Each advertised prefix includes BGP attributes
  - prefix + attributes = “route”
- ◆ 3 important attributes:
  - **AS-PATH**: contains a list of ASes through which prefix advertisement has passed
    - When Router-C receives the announcement: 12.3.6/24: AS4, AS1
  - **NEXT-HOP**: indicates specific internal-AS router to next-hop AS
    - may be multiple links from current AS to next-hop-AS
  - **Local-Preference**: indicates policy preference in path selection
    - Injected into BGP update at border router (by C, D and E)
    - Used by internal routers to decide whether going through AS2 or AS4 to reach 12.3.6/24



## BGP routing policy:

a provider advertises all prefixes to its customer ASes;  
a customer does not advertise prefixes between providers

important



A,B,C are **provider network ASes**

X,W,Y are customer ASes (of provider networks)

X is **dual-homed**: attached to two provider networks


X does not want to forward traffic from B to C

.. so X will not advertise to B any **route** it learned from C

(i.e. IP prefixes with AS path)

important

# BGP routing policy: a provider does not pass prefixes that are not its own to another providers



A advertises to B the path **[1.2.3/24: AW]**

B advertises to X the path **[1.2.3/24: BAW]**

Would B advertise to C the path **[1.2.3/24: CBAW]**?


No way! B gets no “revenue” for routing CBAW since neither W nor C are B’s customers

B wants to route *only* to/from its customers!



# No Valley Policy

important



# Why different Intra- and Inter-AS routing ?

## ◆ Policy:

- Inter-AS: admin wants control over how its traffic routed, who routes through its net.
- Intra-AS: single admin, so no policy decisions needed

## ◆ Scale:

- hierarchical routing saves table size, reduced update traffic

## ◆ Performance:

- Intra-AS: can focus on performance
- Inter-AS: policy may dominate over performance

# A Quick Summary of Internet Routing

- ◆ OSPF: a link-state routing protocol
  - Each router sends Link-State Packet containing
    - ID of the node that created the LSP
    - a list of direct neighbors, with link cost to each
    - sequence number (SEQ) for this LSP message
    - time-to-live (TTL) for information carried in this LSP
  - LSPs are sent periodically, or whenever changes happen
    - flooded everywhere, reliably
  - Neighbor routers use Hello msgs to keep track each other
- ◆ BGP: a path-vector (like distance vector by with paths) routing protocol
  - Running over TCP connection
  - Propagate reachable IP prefixes

# Routing Security

## ◆ Intra-domain routing

- Controlled by a single party, so could be secured using shared secrets
- Still there is an issue if routers are getting compromised

## ◆ Inter-domain routing

- Multiple parties
- Not everyone behaves correctly
  - Configuration errors
  - Malicious activity

# Unintended Behavior in BGP

- ◆ Route hijacking
  - an AS announcing somebody else's prefix(es) as they own
  - depending on AS-PATH, some part of the Internet will prefer hijacked path
    - Google DNS hijacking in 2014 by Turkish ISPs
    - Youtube blackout in 2008 caused by Pakistan ISPs
    - many more
- ◆ Man-in-the-middle (diverting routes)
- ◆ Announcing unused prefixes
  - routes disappear after finishing some communications, e.g., sending spam
- ◆ Using unallocated AS numbers
  - avoiding legal tracing

<http://cyclops.cs.ucla.edu/>